

# Intro to Text Analysis

*Sebastián Vallejo*

*12/5/2019*

## Contents

<b>1</b>	<b>Outline of the workshop</b>	<b>2</b>
1.1	What will you need? . . . . .	2
1.2	What are my sources? . . . . .	2
<b>2</b>	<b>Text as data: Some theory</b>	<b>3</b>
2.1	Why text? . . . . .	3
2.2	Why use the help of a computer? . . . . .	3
2.3	Four principles of automated text analysis (from Grimmer and Stewart (2013)) . . . . .	3
<b>3</b>	<b>The corpus (in theory)</b>	<b>4</b>
3.1	Text as data: Some specifics . . . . .	4
3.2	What counts as a document? . . . . .	4
3.3	Where can I get my hands on some juicy corpora? . . . . .	4
3.4	Some final words on documents . . . . .	4
<b>4</b>	<b>The corpus (in practice)</b>	<b>5</b>
<b>5</b>	<b>Pre-processing the corpus</b>	<b>9</b>
5.1	Before the “bag-of-words” . . . . .	12
<b>6</b>	<b>Analyzing our corpus: Keywords in context and collocations</b>	<b>14</b>
6.1	Keywords in context . . . . .	14
<b>7</b>	<b>Analyzing a corpus: Features</b>	<b>16</b>
7.1	(10) Most frequently used words . . . . .	17
7.2	Lexical diversity . . . . .	18
7.3	TF-IDF . . . . .	19
7.4	Wordclouds . . . . .	20
<b>8</b>	<b>Analyzing a corpus: Dictionary-based approaches</b>	<b>21</b>
8.1	Corpus-specific dictionary . . . . .	21
8.2	LIWC and Sentiment Analysis . . . . .	23
<b>9</b>	<b>Analyzing a corpus: Topic models and document distance</b>	<b>27</b>
9.1	Document-term Matrix . . . . .	27
9.2	Structural Topic Models (STM) . . . . .	29
<b>10</b>	<b>Analyzing a corpus: Scaling models</b>	<b>38</b>
<b>11</b>	<b>Analyzing a corpus: Natural Language Processing (NLP)</b>	<b>42</b>
<b>12</b>	<b>Other techniques not covered</b>	<b>46</b>
12.1	Word networks . . . . .	46
12.2	Cosine-similarity . . . . .	49
12.3	Word embeddings . . . . .	49
<b>13</b>	<b>Goodbye Notes</b>	<b>51</b>

# 1 Outline of the workshop

1. Text as data: Brief overview of the theory behind text analysis
2. The corpus
3. Analyzing our corpora
  - Features
  - Dictionary-based approaches
  - Estimating distance
  - Topic models
  - Natural language processing (NLP)
  - Text networks
  - Neural networks / word embeddings
4. Conclusion and parting remarks

Note: The point of this workshop is not to for you to leave an expert on text analysis, but rather for you to have a taste of what can be achieved (i.e. what substantive questions can be answered) when using text analysis techniques. The code provided can help you get started, but you will need to explore each method more in detail if you want to apply it to your research.

## 1.1 What will you need?

If you want to follow along in your computer, you should have spaCyR intalled. In addition to spaCyR, you should have the following packages installed:

- quanteda
- quanteda.dictionaries
- stm
- tidyverse
- ggplot2
- lme4
- lattice
- ggthemr (to prettify)
- dplyr

Text Analysis is a computer-intensive task. spaCyR and quanteda processes can consume a lot of your RAM, so take that into account when running your code.

## 1.2 What are my sources?

Much of the material/ideas for the workshop were taken from the following sources:

- Kenneth Benoit's website
- The spaCyR GitHub page
- Chris Bail's Course
- Elliot Ash's Course
- Will Lowe's Text Analysis workshop at IQMR
- The Internet

If you are interested in applied and/or theoretical readings on text analysis, here is a short list to get you started:

- Grimmer and Stewart (2013) - Text as Data: The Promise and Pitfalls of Automatic Content Analysis Methods for Political Texts **in Political Analysis**
- Lucas et al. (2015) - Computer-Assisted Text Analysis for Comparative Politics **in Political Analysis**

- Blei (2012) - Probabilistic Topic Models **in Communications of the ACM**
- **Poetics**, Volume 41, Special Issue on “Topic Models and the Cultural Sciences”
- Slapin and Proksch (2008) - A Scaling Model for Estimating Time-Series Party Positions from Texts **in AJPS**
- Welbers et al. (2017) - Text Analysis in R **in Communication Methods and Measures**

Let’s start!

## 2 Text as data: Some theory

### 2.1 Why text?

1. Politicians *love* to speak.
2. Bureaucrat *love* to write.
3. Machirulos *love* to tweet.

There is text in every aspect of politics: debates on legislation, peace treaties, news reports, political manifestos, campaign speeches, social media, etc. Not only is text ubiquitous, but it is produced at the time (sometimes in real time).

### 2.2 Why use the help of a computer?

1. Humans are great at understanding and analyzing the content of texts. Computers are not.
2. Humans are also great at not being able to read thousands of documents in minutes, organize that text, classify that text, scale that text, and then produce pretty graphics from that text. Computers are not not that. (I stand by the double negative)

We will be learning about the latter.

Note: Text analysis is not a field, it is a tool. Think about it as a regression where the data are words instead of numbers. Thus, the fanciest of text analysis techniques is no good without a well thought out, substantive, theoretically motivated question.

### 2.3 Four principles of automated text analysis (from Grimmer and Stewart (2013))

1. All Quantitative Models of Language Are Wrong—But Some Are Useful.
  - Data generation process for text is unknown.
  - Language is too complex for computers to correctly decipher (e.g. “Time flies like an arrow. Fruit flies like a banana.”).
  - Since language is so context-specific, more complex models are rarely more useful for the analysis of texts.
2. Quantitative Methods Augment Humans, Not Replace Them
  - You need to read the text, know the text, be the text.
  - Computers organize, direct, and suggest.
  - Humans read and interpret.
3. There Is No Globally Best Method for Automated Text Analysis
  - Different needs require different methods.
  - Even when you have the same needs, the same model might not fit the data.

#### 4. Validate, Validate, Validate

- Outputs can be misleading (or simply wrong).
- It is incumbent upon you, the researcher, to validate the use of automated text analysis.

Now we can start with some (sample) text analysis. Note that we will be doing basic analysis (as generalizable as it gets) of text, and then use some models as examples. To see the extent of what can be done with different models, I suggest you read Grimmer and Stewart (2013).

### 3 The corpus (in theory)

To analyze text, we first need a corpus, a large and structured set of texts for analysis. The units (texts) of the “structured set of texts” can be anything you want them to be: a complete speech, each paragraph of the speech, or each sentence within that paragraph. The relevance of the unit will depend on your research question.

#### 3.1 Text as data: Some specifics

Text data is a sequence of characters called **documents**. The set of documents is the **corpus**.

- Text data is **unstructured**
- There is information we want, mixed in with (A LOT) of information we do not.
- We need to separate the wheat from the chaff.

Note: All text analysis methods will produce some information. The *art* lies on the ability of the researcher to figure out what is valuable and what is not.

#### 3.2 What counts as a document?

The unit of analysis when using text as data will depend on the question you are asking. For example:

- If you are looking at how politicians react to different types of economic crises, then each document produced after a crisis would be the unit of analysis.
- If you are looking at how politicians differ within a campaign, then you might aggregate all the texts produced by a candidates during a campaign as the unit of analysis.
- If you are looking at how politicians address different topics within a campaign, then your unit of analysis might be a section or paragraph of a manifesto.

#### 3.3 Where can I get my hands on some juicy corpora?

1. Chris Bail curates a list of corpora already compiled and ready to use.
2. Governments produce text ALL THE TIME. It is, usually, publicly available and, depending on the country/organization, easily accessible.
3. Scrape the webs. Websites can make it difficult to scrape data with restrictive terms of use (i.e. bot-blockers or, worse, javascript). There are creative ways to get around these. (If you are interested in web-scraping, let me know and I can help you get started with some code.)

#### 3.4 Some final words on documents

Original corpora are rarely ready to be analyzed without some previous cleaning. You often want to get rid of hyphenations at line brakes, table of contents, indexes, etc. All of these are corpus-specific and require attention ahead of time.

- Learn how to use regular expressions (regex). The `stringr` package in R or the Python package `re` are useful tools.
- While useful, regex is tedious to learn. Check out Sanchez 2013 for a good guide.

Some data are only available as non-searchable PDFs or images. These need to be converted to text before R (or Python) can read them. I use ABBYY FineReader, which is ‘expensive’ but might be available at your (old/next) university library. Joe Sutherland has an open-source OCR (Optical Character Recognition).

Finally, I advise against using spell checkers. Most corpora use specialized language that would be flagged by standard spell-checkers (and I have not found one that can be automated to check text in Spanish). In most empirical contexts, we can safely assume that spelling errors (especially OCR errors) are uncorrelated with treatment assignment.

(Ask me about some of the other strengths weaknesses of using data from social media.)

## 4 The corpus (in practice)

We need texts for text analysis. Luckily, we have a dataset of 5640 tweets from 3885 unique users containing the word ‘capitalism’. The dataset has been tinkered with and cleaned and is ready to be processed.

Before we start, let’s load the required packages. Remember that `spacyR` is a Python wrapper that needs to first be initialized.

```
rm(list=ls(all=TRUE))
library(tidyverse)
library(spacyr)
library(stm)
library(quanteda)
library(quanteda.dictionaries)
library(dplyr)
library(lattice)
library(ggplot2)
library(ggrepel)
library(lme4)
library(ggthemr)
ggthemr("fresh")
spacy_initialize()
```

Let’s load our dataset and see how it looks.

```
load("data_capitalism.Rdata")
data_capitalism <- data_capitalism[!data_capitalism$text_clean=="",]
data_capitalism %>% glimpse()

## Observations: 5,627
## Variables: 13
## $ text          <chr> "if you're having trouble understanding capitalism j...
## $ friendsRT     <int> 637, 576, 674, 116, 2937, 7721, 11768, 3011, 2838, 1...
## $ followersRT   <int> 1126, 3407, 36590, 5080, 6088, 12074, 15690, 3277, 1...
## $ timeRT        <dbl> 18195, 18228, 18223, 18230, 18226, 18222, 18223, 182...
## $ nameauth      <chr> "shnupz", "aniceburrito", "CaseyExplosion", "HairyBo...
## $ likeRT        <int> 45195, 4349, 2850, 57, 78, 183, 34, 110, 3750, 133, ...
## $ retweetRT     <int> 7258, 590, 1031, 5, 16, 41, 16, 44, 1113, 55, 12, 90...
## $ to_membership <dbl> 3, 3, 3, 3, 6, 6, 6, 3, 6, 6, 6, 6, 6, 6, 6, 9, 9...
## $ to_ind        <dbl> 82, 115, 238, 3, 11, 17, 12, 280, 282, 16, 8, 98, 98...
## $ mem_name      <chr> "POC Anti-Capitalism", "POC Anti-Capitalism", "POC A..."
```

```
## $ mem_name_dummy <chr> "Anti-Capitalism", "Anti-Capitalism", "Anti-Capitali...
## $ text_clean <chr> "if you're having trouble understanding capitalism j...
## $ date_created <date> 2019-10-26, 2019-11-28, 2019-11-23, 2019-11-30, 201...
```

To analyze the text found in the dataset, we need to create a `corpus` object. The `corpus()` function (from the `quanteda` package) does just this. The main object has to be a character vector. (The `readtext()` function from the `readtext` package can read text-formatted files).

```
cap_corp <- corpus(data_capitalism$text_clean,
  docvars = data.frame(author = data_capitalism$nameauth,
    time = data_capitalism$timeRT,
    in_degree = data_capitalism$to_ind,
    followers = data_capitalism$followersRT,
    RT_count = data_capitalism$retweetRT,
    left_right = data_capitalism$mem_name,
    left_right_dum = data_capitalism$mem_name_dummy),
  metacorpus = list(source = "Twitter",
    notes = "Scrapped on Nov. 30 - Dec. 3, 2019"))

summary(cap_corp, 10)
```

```
## Corpus consisting of 5627 documents, showing 10 documents:
```

```
##
##   Text Types Tokens Sentences      author   time in_degree followers
##   text1     22    23         1      shnupz 18195      82      1126
##   text2     13    14         2    aniceburrito 18228     115      3407
##   text3     44    53         2 CaseyExplosion 18223     238     36590
##   text4      8     8         1 HairyBoomerDude 18230      3      5080
##   text5     36    51         4 brianjtrautman 18226     11      6088
##   text6     10    10         1      LostDiva 18222     17     12074
##   text7     39    55         3      safeagain1 18223     12     15690
##   text8     22    29         2      lvl28mage 18229    280      3277
##   text9      7     7         1   davidsirota 18226    282     153820
##   text10    17    20         1      jayasax 18222     16     18902
##   RT_count      left_right left_right_dum
##   7258    POC Anti-Capitalism Anti-Capitalism
##   590     POC Anti-Capitalism Anti-Capitalism
##   1031    POC Anti-Capitalism Anti-Capitalism
##   5       POC Anti-Capitalism Anti-Capitalism
##   16 White Anti-Capitalism Anti-Capitalism
##   41 White Anti-Capitalism Anti-Capitalism
##   16 White Anti-Capitalism Anti-Capitalism
##   44     POC Anti-Capitalism Anti-Capitalism
##   1113 White Anti-Capitalism Anti-Capitalism
##   55 White Anti-Capitalism Anti-Capitalism
##
## Source: Twitter
## Created: Sun Dec 8 20:19:46 2019
## Notes: Scrapped on Nov. 30 - Dec. 3, 2019
```

A corpus object works similar to a normal dataset, in that each document (in our case, each tweet) is an observation that has additional covariates (i.e. *docvars*) that describe it.

```
cap_corp[c(1:5)]
```

```
##
```

```
##
##
##
##
##      "\"Fascism is capitalism in decay.\"\\n\\nI didn't even understand what th
##
##
##
## "@davidsirota No stage - it's just capitalism. Stages aren't distinguishable with capitalism. Only i
```

You can explore a corpus as you would explore other lists (with brackets). To get all the texts you can `texts(cap_corp)`, but you don't want to do that. We can subset a corpus. Here are all the tweets from TW authorities with more than 90K followers (and less than 200 words... I know, useless, but you might appreciate the code to do it):

```
cap_subset <- corpus_subset(cap_corp,
                             followers > 90000 & ntype(cap_corp) < 200)
summary(cap_subset, 15)
```

```
## Corpus consisting of 424 documents, showing 15 documents:
```

```
##
##      Text Types Tokens Sentences      author  time in_degree followers
##      text9      7      7          1      davidsirota 18226      282      153820
##      text11     16     16          2      mikebarnicle 18225        8      133427
##      text12     25     52          4      peterdaou 18231       98      281684
##      text13     27     28          2      peterdaou 18232       98      281675
##      text17     14     16          3 RedNationRising 18229      118      170751
##      text18     24     25          1 KurtSchlichter 18232      131      211589
##      text19     33     46          3 LacyJohnsonMN 18229      258      103051
##      text20     39     48          4      ACTBrigitte 18227      284      193371
##      text25     29     33          2 BreitbartNews 18232      350     1207329
##      text27     19     19          1 BridgetPhetasy 18231      756      153876
##      text54     40     47          3      iheartmindy 18231      362      161898
##      text55     21     25          2      thecjpearson 18229       90      320951
##      text56     20     22          2      NikkiHaley 18225      294      524277
##      text60     24     25          2 SandraSentinel 18229       59      169095
##      text61     46     57          5      mikandynothem 18232      168      220387
```

```
##      RT_count      left_right left_right_dum
##      1113 White Anti-Capitalism Anti-Capitalism
##      12 White Anti-Capitalism Anti-Capitalism
##      90 White Anti-Capitalism Anti-Capitalism
##      355 White Anti-Capitalism Anti-Capitalism
##      717      Capitalism      Capitalism
##      528      Capitalism      Capitalism
##      1352      Capitalism      Capitalism
##      2074      Capitalism      Capitalism
##      1727      Capitalism      Capitalism
##      4115      Capitalism      Capitalism
##      1093      Capitalism      Capitalism
##      61      Capitalism      Capitalism
##      2741      Capitalism      Capitalism
##      24      Capitalism      Capitalism
##      554      Capitalism      Capitalism
```

```
## Source: Twitter
```

```
## Created: Sun Dec 8 20:19:46 2019
## Notes: Scrapped on Nov. 30 - Dec. 3, 2019
```

We might be interested in sentences rather than complete tweets.

```
cap_sentences <- corpus_reshape(cap_corp, to = "sentences") # or "paragraphs"
summary(cap_sentences,8)
```

```
## Corpus consisting of 12172 documents, showing 8 documents:
```

```
##
##      Text Types Tokens Sentences      author  time in_degree followers
## text1.1      22      23          1      shnupz 18195          82        1126
## text2.1      12      13          1    aniceburrito 18228          115        3407
## text2.2       1       1          1    aniceburrito 18228          115        3407
## text3.1       7       8          1 CaseyExplosion 18223          238       36590
## text3.2      39      45          1 CaseyExplosion 18223          238       36590
## text4.1       8       8          1 HairyBoomerDude 18230           3        5080
## text5.1       8       8          1 brianjtrautman 18226           11       6088
## text5.2       6       6          1 brianjtrautman 18226           11       6088
## RT_count      left_right left_right_dum
##      7258 POC Anti-Capitalism Anti-Capitalism
##      590 POC Anti-Capitalism Anti-Capitalism
##      590 POC Anti-Capitalism Anti-Capitalism
##     1031 POC Anti-Capitalism Anti-Capitalism
##     1031 POC Anti-Capitalism Anti-Capitalism
##       5 POC Anti-Capitalism Anti-Capitalism
##      16 White Anti-Capitalism Anti-Capitalism
##      16 White Anti-Capitalism Anti-Capitalism
##
## Source: Twitter
## Created: Sun Dec 8 20:19:47 2019
## Notes: corpus_reshape.corpus(cap_corp, to = "sentences")
```

```
cap_sentences[1:5]
```

```
##
##
##
##
##
##
##
##
## "I didn't even understand what that meant a few years back, but the wealthy funnelling money to white
```

Since these are tweets and people often tweet one sentence at a time, this conversion might be moot for our corpus. But in longer texts narrowing down the unit of analysis can be helpful, especially if we are trying to estimate topic models (more on this later).

The reshape function can divide texts into “sentences” and “paragraphs”. `corpus_reshape()` uses punctuation marks (e.g. “\n”, “\n”) to determine cuts.



## 5 Pre-processing the corpus

As previously mentioned, our corpora have the information we want, and a lot of information we do not. Uninformative data add noise and reduce the precision of resulting estimates (and are computationally costly). We aim to have a “bag-of-words”, or to convert a corpus  $D$  to a matrix  $X$ . In the “bag-of-words” representation, a row of  $X$  is just the frequency distribution over words in the document corresponding to that row.

Before we do that, we will get rid of all the unwanted information. First, we turn all words to lower-case and get rid of all punctuation and numbers. The `tokens()` command will do this and separate all the texts into tokens.

```
cap_toks <- tokens(cap_corp,
                  remove_numbers = T,
                  remove_punct = T,
                  remove_separators = TRUE,
                  remove_twitter = T,
                  remove_hyphens = T,
                  remove_url = T)
```

**Tokens** (`ntoken`) is a fancy name for “word”. These contain all the information we need to run our models. Yet, there is still a lot of noise.

```
cap_coll <- textstat_collocations(cap_toks)
head(cap_coll, 20)
```

##	collocation	count	count_nested	length	lambda	z
## 1	capitalism is	789	0	2	2.362610	54.56040
## 2	of capitalism	686	0	2	1.818068	41.17951
## 3	free market	88	0	2	6.805751	40.79497
## 4	this is	285	0	2	2.898294	40.58134
## 5	to be	264	0	2	2.877196	37.78918
## 6	climate change	75	0	2	8.099856	37.64941
## 7	if you	145	0	2	3.750860	37.62064
## 8	we have	130	0	2	3.619659	35.53517
## 9	working class	57	0	2	6.626296	35.23928
## 10	in the	450	0	2	1.912162	34.93206
## 11	late stage	57	0	2	8.194708	34.43979
## 12	we need	92	0	2	4.611996	34.13644
## 13	under capitalism	196	0	2	3.900645	33.38231
## 14	mark ruffalo	78	0	2	9.382803	32.39438
## 15	is a	362	0	2	1.881855	32.06737
## 16	economic revolution	47	0	2	6.941818	32.06536
## 17	people who	89	0	2	3.863746	31.64427
## 18	those who	59	0	2	5.106866	31.18187
## 19	should be	65	0	2	4.770880	30.16049
## 20	has been	57	0	2	4.918624	29.77172

(Quick detour: Collocations bundle together a set number of words -also known as ngrams- that appear next to each other. The default is 2, but we can set it at any length we want.)

Some words are “useless”. Let’s get rid of the *stopwords*. (For a take on when stopwords are informative, check Pennebaker (2011).)

```
cap_toks_stop <- tokens_remove(cap_toks,
                               stopwords(language = "en"),
                               padding = F)
```

```
cap_toks_stop <- tokens_remove(cap_toks_stop, "amp")
cap_toks_stop <- tokens_remove(cap_toks_stop, "capitalism")
cap_coll <- textstat_collocations(cap_toks_stop)
head(cap_coll, 20)
```

##	collocation	count	count_nested	length	lambda	z
## 1	free market	88	0	2	6.176197	36.84064
## 2	climate change	75	0	2	7.462913	34.54893
## 3	working class	57	0	2	5.979366	31.75400
## 4	late stage	57	0	2	7.748096	31.36852
## 5	mark ruffalo	78	0	2	8.744120	30.13495
## 6	economic revolution	47	0	2	6.304274	29.00872
## 7	killing us	52	0	2	5.755766	25.98845
## 8	ethical consumption	35	0	2	8.636356	25.37620
## 9	means production	26	0	2	6.504765	23.97048
## 10	ruffalo calls	26	0	2	7.036844	23.82651
## 11	bernie sanders	30	0	2	8.042982	23.30022
## 12	calls economic	31	0	2	6.596260	23.21391
## 13	social media	27	0	2	5.845772	23.12489
## 14	economic system	39	0	2	4.323610	22.97572
## 15	everyone else	24	0	2	5.697881	22.19599
## 16	climate crisis	24	0	2	5.528306	21.88571
## 17	worth million	18	0	2	6.113556	20.61271
## 18	gift buying	16	0	2	7.764324	20.54672
## 19	white supremacy	37	0	2	8.370088	20.29066
## 20	years ago	25	0	2	8.068831	20.26880

Finally, we might want to stem our tokens.

```
cap_toks_stem <- tokens_wordstem(cap_toks_stop)
cap_coll <- textstat_collocations(cap_toks_stem)
head(cap_coll, 20)
```

##	collocation	count	count_nested	length	lambda	z
## 1	free market	110	0	2	6.018092	40.29298
## 2	climat chang	74	0	2	7.027756	35.06895
## 3	late stage	57	0	2	7.704935	31.43510
## 4	mark ruffalo	86	0	2	8.721140	31.33799
## 5	econom revolut	47	0	2	6.072575	28.60968
## 6	kill us	58	0	2	4.635673	27.88024
## 7	work class	59	0	2	4.353081	26.88487
## 8	ethic consumpt	34	0	2	8.028629	25.64903
## 9	berni sander	30	0	2	7.928019	23.15127
## 10	call econom	34	0	2	4.597643	22.79844
## 11	econom system	40	0	2	4.051867	22.18342
## 12	everyon els	24	0	2	5.605248	21.96526
## 13	climat crisi	24	0	2	5.536514	21.90668
## 14	rule class	25	0	2	5.482741	21.65451
## 15	ruffalo call	26	0	2	5.105934	21.54630
## 16	mean product	26	0	2	4.889848	21.19389
## 17	white supremaci	37	0	2	8.224788	20.93996
## 18	net worth	22	0	2	8.152715	20.46240
## 19	leecamp sensand	24	0	2	10.059114	19.97046
## 20	black friday	55	0	2	8.993914	19.66205

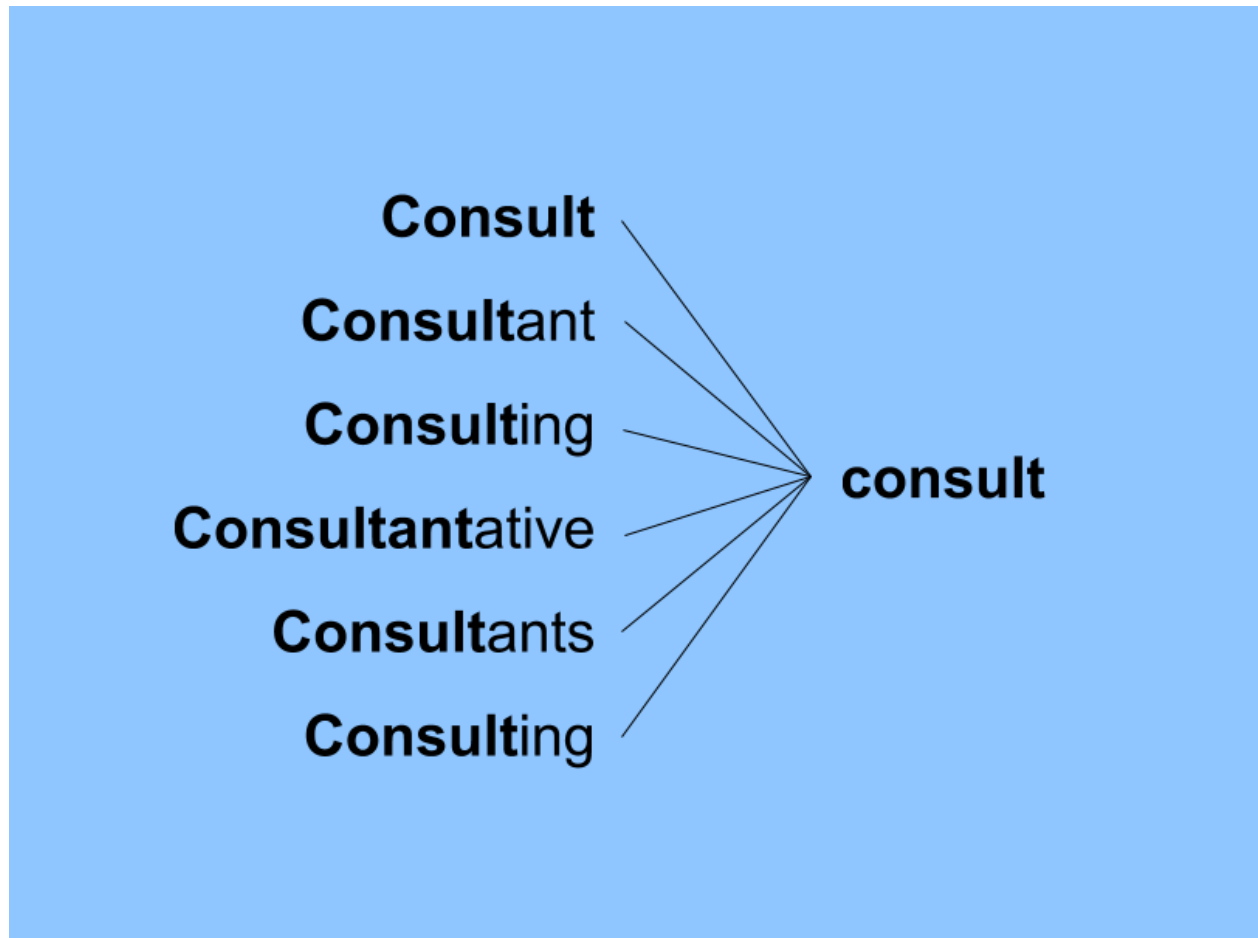


Figure 1: Example of stemming

## 5.1 Before the “bag-of-words”

We have *just* pre-processed the data, but the number of documents (e.g. tweets) and the (pre-processed) length of these documents already provide an interesting set of variables for analysis.

For example: - How do major capitalist events affect the production of tweets from capitalists and anti-capitalists? - What is the relation between in-degree and effort and quality of tweets?


*# How do major capitalist events affect the production of tweets from capitalists and anti-capitalists?*

```
data_capitalism <- data_capitalism %>%
  group_by(date_created, mem_name_dummy) %>%
  mutate(count_tweets = n())

data_capitalism_res <- data_capitalism[!duplicated(data_capitalism[c(11,13)]),]
data_capitalism_res <- data_capitalism_res[data_capitalism_res$date_created > "2019-11-19",]

ggplot(data_capitalism_res, aes(x=date_created , y = count_tweets , color = mem_name_dummy)) +
  stat_smooth() +
  scale_color_discrete(name = "") +
  labs(x = "Date", y = "Count of Tweets") +
  theme(axis.text.x = element_text(angle = 90, hjust = 1),
        legend.position="bottom") +
  geom_vline(xintercept = as.Date("2019-11-29"), color = "black", linetype = "dashed")

## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



the Data I-1.bb

Anti-Capitalism Capitalism

*# What is the relation between in-degree and effort of tweets?*


```

data_capitalism$length_tweet <- ntoken(char_tolower(data_capitalism$text_clean),
                                       remove_punct = TRUE)

data_capitalism$length_type <- ntoken(char_tolower(data_capitalism$text_clean),
                                       remove_punct = TRUE) #repeated words are cut..

ggplot(data_capitalism, aes(x=to_ind , y = length_tweet)) +
  geom_point() +
  geom_smooth(method = "lm", se = T)+
  labs(x = "In-Degree", y = "Length of Tweets")

```




the Data I-2.bb

```

ggplot(data_capitalism, aes(x=to_ind , y = length_type)) +
  geom_point() +
  geom_smooth(method = "lm", se = T)+
  labs(x = "In-Degree", y = "Number of Types per Tweets")

```



the Data I-3.bb

## 6 Analyzing our corpus: Keywords in context and collocations

Finally, we are ready to analyze our corpus. We will start by the basic: keywords in context.

### 6.1 Keywords in context

Say we are interested in the way capitalist and anti-capitalist describe **capitalism**. Let's see how they talk about **capitalism** using the “keyword-in-context” function:

```
cap_talk <- kwic(cap_toks_stop, "capital*", window=10) # 10 words before and after the word in question
head(cap_talk, 20)
```

```
##
## [text35, 3]
## [text44, 20]
## [text64, 3]
## [text64, 16]
## [text69, 17]
## [text101, 1]
## [text102, 18]
## [text111, 12]
## [text125, 7]
## [text131, 24]
## [text136, 3]
## [text136, 14]
```

```

## [text137, 18]
## [text150, 1]
## [text151, 1]
## [text159, 6]
## [text160, 1]
## [text164, 2]
## [text168, 17]
## [text171, 11]
##
##
## know Fascism
## services socialism reality whatever programs creates always constantly threat ruling
## Every time
## calls population control matter dress rank pseudo science used obscure
## financial weight stop Luthor Corp various drug lords parents sycophants
##
## Finance rules benefit elite wealthy leave masses behind struggling wake
## systematically designed extremely powerful counter revolutionary force many ways global
## Unpopular opinion think qualities fetishize entrepreneurship
## coup attempt Panama weeks later US invaded overthrew gov CIA
## Arguing defence
## ongoing effort suck blood poor giving junk return argument privilege
## leaguer Rhodes Scholar gone work anywhere world chose go work
##
##
## Lech Walesa Harvard deplore US
##
## Among
## swerves accelerates retreats depending routes open tolls lower one get
## wrote biggest difference Bernie Sanders Elizabeth Warren views constitutes corruption
##
## | Capitalism's |
## | capitalist |
## | Capital |
## | capitalism's |
## | capitalist |
## | Capitalism's |
## | Capitalism's |
## | capitalism's |
## | capitalist |
## | Capitalism's |
## | capitalism's |
## | capitalism's |
## | capitalism's |
## | capitalism's |
## | Capitalism's |
## | Capitalism's |
## | capitalism's |
## | Capitalism's |
## | capitalism's |
## | Capitalism's |
## | capitalism's |
## | capital |
## | capitalist |
##
## life support system decay activate stomp dissidents alternatives recoils back
## class
## approaches crisis calls population control matter dress rank pseudo science

```

```
## role ecocide set scientific precedent genocide
## values Comrade Wayne enough sabotoge efforts thwart x
## Failures Ignited Worldwide Protests
## inability answer social economic needs people
## crown jewel really need stop looking liberatory art constantly disappointed
## success behaviors prop rape culture example saw book business section
## Invisible Army
## ongoing effort suck blood poor giving junk return argument privilege
## bidding Poor people deserve nice things value money just like
## despicable mercenaries McKinsey co call sell call zealot
## Housing Crisis Socialist Alternative Democrat Party forces America sees calling
## Failures Ignited Worldwide Protests TheRealNews
## global decline course decline good empire Self congratulatory delusions
## self destruction US economy's future depends increasingly quality quantity educated
## basic injustices small minority employers can hire fire majority employees
## profitable destination faster One woman's corruption another woman's
## economies Let know think
```

If we put together all these words in one whole text, we can see how they are related in that context.

```
cap_talk_context <- paste(cap_talk$pre, cap_talk$post, collapse = " ")
cap_talk_coll <- textstat_collocations(cap_talk_context, size = 3)
head(cap_talk_coll, 20)
```

##		collocation	count	count_nested	length	lambda	z
## 1		production means means	2	0	3	3.344771	1.3068288
## 2		economic needs people	2	0	3	3.278764	1.2789896
## 3		know people pay	2	0	3	2.825768	1.0993908
## 4		capitalist also actually	2	0	3	2.329066	1.0192045
## 5		capital can make	2	0	3	2.170820	0.9495177
## 6		needs instead just	2	0	3	2.359568	0.9151489
## 7		paid work get	2	0	3	2.343121	0.9094750
## 8		ignited protests worldwide	6	0	3	2.211441	0.8701963
## 9		labor capitalist mode	3	0	3	1.743161	0.7611189
## 10		rich labor rich	3	0	3	1.946246	0.7581333
## 11		ignited worldwide protests	8	0	3	1.794087	0.7034010
## 12		anti capitalist also	2	0	3	1.289242	0.6807563
## 13		labor nature work	2	0	3	1.540523	0.6678149
## 14		social economic needs	2	0	3	1.456625	0.6348281
## 15		instead just accumulation	2	0	3	1.620961	0.6242451
## 16		nature work social	2	0	3	1.339011	0.5800195
## 17		democracy pro democratic	2	0	3	1.378495	0.5324733
## 18		labor consuming make	3	0	3	1.358459	0.5208067
## 19		society poor people	2	0	3	1.038841	0.4635130
## 20		hey know people	2	0	3	1.197946	0.4550775

As a reference, the  $\lambda$  score is a measure of the times  $K$  specific consecutive tokens happen (in our second example,  $K = 3$ ), given all the  $K$  consecutive tokens possibilities.

## 7 Analyzing a corpus: Features

The features of a corpus can give us clues about the characteristics of our text. At this point, we are already treating our documents as “bags-of-words”. Having tokenized our corpus means that we are interested in each word, how often they appear, in conjunction with what other words, etc. We are going to create a



Document-Feature Matrix (*dfm*) object. I will going into detail of what is a *dfm*. For now, just think of it a dataset where each row if the count of words of each document.

```
dfm_toks <- dfm(cap_toks_stop)
```

Note that my *dfm* object contains all the document level variables I added to my corpus at the beginning.

## 7.1 (10) Most frequently used words


Simply:

```
cap_freq <- textstat_frequency(dfm_toks, n = 5, groups = "left_right")
head(cap_freq, 15)
```

##	feature	frequency	rank	docfreq	group
## 1	socialism	247	1	217	Capitalism
## 2	markruffalo	214	2	206	Capitalism
## 3	people	178	3	156	Capitalism
## 4	us	163	4	138	Capitalism
## 5	mark	139	5	128	Capitalism
## 6	people	348	1	305	POC Anti-Capitalism
## 7	like	254	2	223	POC Anti-Capitalism
## 8	just	253	3	238	POC Anti-Capitalism
## 9	us	221	4	185	POC Anti-Capitalism
## 10	can	196	5	186	POC Anti-Capitalism
## 11	people	298	1	255	White Anti-Capitalism
## 12	like	211	2	182	White Anti-Capitalism
## 13	us	208	3	182	White Anti-Capitalism
## 14	just	194	4	178	White Anti-Capitalism
## 15	can	181	5	163	White Anti-Capitalism

Or a plot of the most frequently used words by group:

```
dfm_toks %>%
  textstat_frequency(n = 10, groups = "left_right") %>%
  ggplot(aes(x = reorder(feature, frequency), y = frequency, color = group)) +
  geom_point() +
  coord_flip() +
  labs(x = NULL, y = "Frequency")
```



## 7.2 Lexical diversity

We might be interested in the breadth and variety of vocabulary used in a document. *Lexical diversity* is widely believed to be an important parameter to rate a document in terms of textual richness and effectiveness. Knowing Marxists, we might expect them to use more complex language than capitalists.

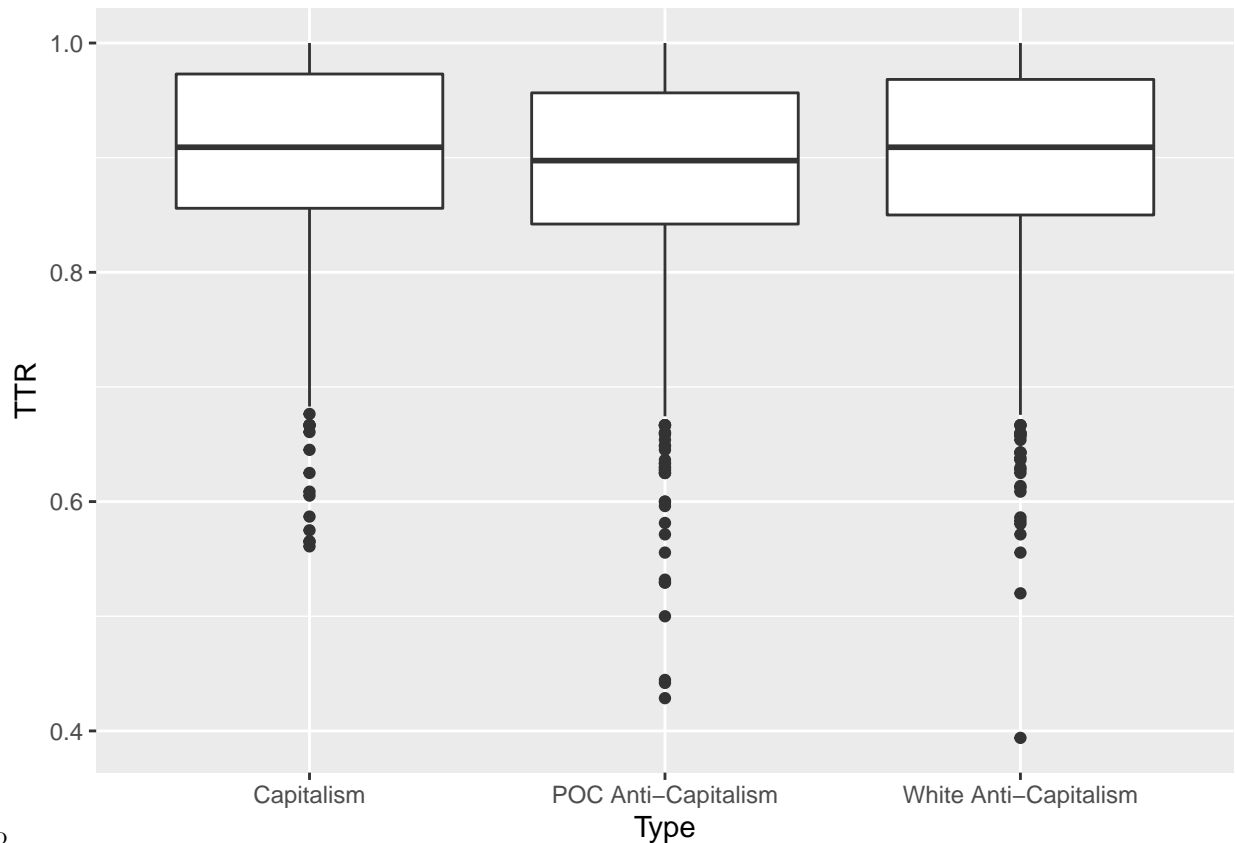
```
dfm_toks_2 <- dfm(cap_toks)
```

```
cap_lexdiv <- textstat_lexdiv(dfm_toks_2, measure = "TTR")
tail(cap_lexdiv, 5) # Does not tell us much, but...
```

```
##      document      TTR
## 5623 text5623 0.9473684
## 5624 text5624 0.8571429
## 5625 text5625 0.9729730
## 5626 text5626 0.7804878
## 5627 text5627 0.9743590
```

```
to_plot <- cbind.data.frame(data_capitalism, cap_lexdiv)
```

```
ggplot(to_plot, aes(x= factor(mem_name), y=TTR)) +
  geom_boxplot() +
  labs(x = "Type")
```



There are many measures of lexical diversity, each measuring something slightly different. `textstat_lexdiv()` includes many and you can check which adapts best to your needs here.

### 7.3 TF-IDF

Looking at simple frequencies might hide some important document features. As in everything in the social sciences, we can always complicate it a bit more. Enter TF-IDF: “Term-frequency / Inverse-document-frequency”. TF-IDF weighting up-weights relatively rare words that do not appear in all documents. Using term frequency and inverse document frequency allows us to find words that are characteristic for one document within a collection of documents.

```
head(dfm_toks[, 5:15])
```

```
## Document-feature matrix of: 6 documents, 11 features (83.3% sparse).
## 6 x 11 sparse Matrix of class "dfm"
##      features
## docs  websites use hate cronyism poster really pushes top beautiful fascism
## text1      1  1  1      0      0      0      0  0      0      0
## text2      0  0  0      1  1      1      1  1      1      0
## text3      0  0  0      0  0      0      0  0      0      1
## text4      0  0  0      0  0      0      0  0      0      0
## text5      0  0  0      0  0      0      0  0      0      0
## text6      0  0  0      0  0      0      0  0      0      0
##      features
## docs  decay
## text1     0
## text2     0
```

```
## text3 1
## text4 0
## text5 0
## text6 0
```

```
head(dfm_tfidf(dfm_toks)[, 5:15])
```

```
## Document-feature matrix of: 6 documents, 11 features (83.3% sparse).
## 6 x 11 sparse Matrix of class "dfm"
##           features
## docs  websites      use      hate cronyism  poster  really  pushes    top
## text1 3.449247 1.869463 1.701059 0          0          0          0          0
## text2 0          0          0      2.231763 2.495004 1.568433 3.273156 2.258915
## text3 0          0          0          0          0          0          0          0
## text4 0          0          0          0          0          0          0          0
## text5 0          0          0          0          0          0          0          0
## text6 0          0          0          0          0          0          0          0
##           features
## docs beautiful  fascism    decay
## text1 0          0          0
## text2 2.847187 0          0
## text3 0          1.858182 2.708884
## text4 0          0          0
## text5 0          0          0
## text6 0          0          0
```

If we are building a dictionary, for example, we might want to include words with high TF-IDF values. Another way to think about TF-IDF is in terms of predictive power. Words that are common to all documents do not have any predictive power and receive a TD-IDF value of 0. Words that appear, but only in relatively few document, have greater predictive power and receive a TD-IDF  $> 0$ .

## 7.4 Wordclouds

Wordclouds are silly, but people seem to **love** them. Begrudgingly, I include the code:

```
# In order to group the wordcloud I will create a new dfm with only one group:
dfm_toks_wordcloud <- dfm(cap_toks_stop,
                           groups = "left_right")

# comparison = T divides the word cloud into groups:
textplot_wordcloud(dfm_toks_wordcloud, comparison = T, max_words = 300)
```




```
dict_women <- dictionary(list(mujeres = c("woman", "women", "girl*", "female*")))

cap_women <- liwcalike(data_capitalism$text_clean,
                      dictionary = dict_women)

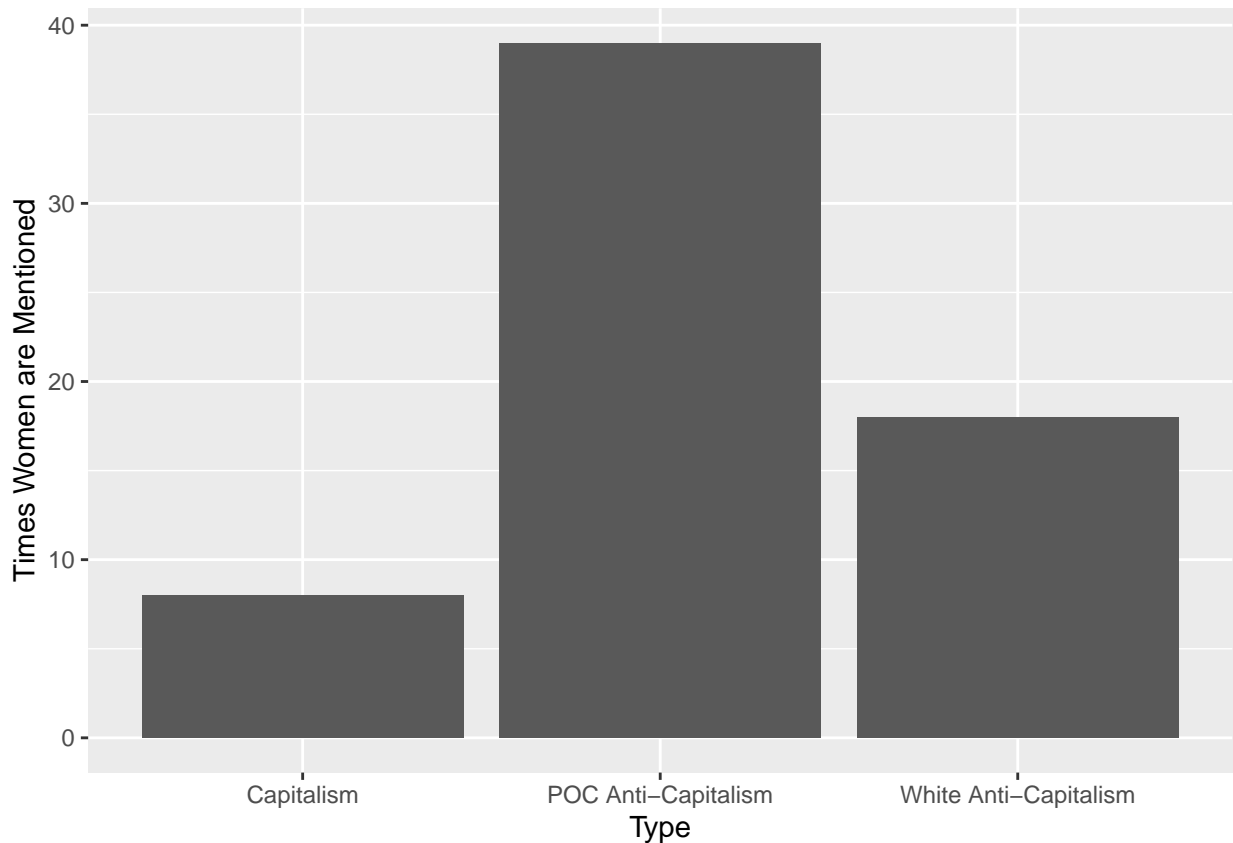
to_plot <- cbind.data.frame(cap_women, data_capitalism)
to_plot$women_dum <- 0
to_plot$women_dum[to_plot$mujeres > 0] <- 1

ggplot(to_plot, aes(x=mem_name, y = mujeres)) +
  geom_boxplot() +
  labs(x = "Type")
```



I-1.bb

```
ggplot(to_plot, aes(x=mem_name, y = women_dum)) +
  geom_bar(stat = "identity") +
  labs(x = "Type", y = "Times Women are Mentioned" )
```



I-2.bb

“The patriarchy is effing vast...”

## 8.2 LIWC and Sentiment Analysis

Let’s do some simple “sentiment analysis” using two dictionaries: GI and NRC.

```
cap_sentimentNRC <- liwcalike(cap_corp,
                             dictionary = data_dictionary_NRC)

cap_sentimentGI <- liwcalike(cap_corp,
                             dictionary = data_dictionary_geninqposneg)

head(cap_sentimentNRC, 15)
```

##	docname	Segment	WC	WPS	Sixltr	Dic	anger	anticipation	disgust	fear
## 1	text1	1	23	23.000000	21.74	30.43	4.35	0.00	4.35	4.35
## 2	text2	2	14	5.500000	21.43	35.71	0.00	7.14	0.00	0.00
## 3	text3	3	53	22.500000	16.98	37.74	1.89	3.77	1.89	3.77
## 4	text4	4	8	7.000000	12.50	0.00	0.00	0.00	0.00	0.00
## 5	text5	5	51	9.250000	27.45	41.18	9.80	0.00	5.88	5.88
## 6	text6	6	10	9.000000	30.00	0.00	0.00	0.00	0.00	0.00
## 7	text7	7	55	14.333333	20.00	7.27	0.00	0.00	0.00	0.00
## 8	text8	8	29	12.000000	34.48	55.17	10.34	0.00	13.79	6.90
## 9	text9	9	7	6.000000	14.29	0.00	0.00	0.00	0.00	0.00
## 10	text10	10	20	19.000000	15.00	5.00	0.00	0.00	0.00	0.00
## 11	text11	11	16	7.000000	18.75	25.00	0.00	6.25	0.00	0.00
## 12	text12	12	52	9.000000	23.08	19.23	1.92	0.00	0.00	3.85

```
## 13 text13      13 28 12.500000 25.00 25.00 0.00      0.00      0.00 7.14
## 14 text14      14 19 5.333333 31.58 68.42 10.53      5.26      5.26 10.53
## 15 text15      15 48 12.000000 33.33 20.83 0.00      0.00      2.08 4.17
##      joy negative positive sadness surprise trust AllPunc Period Comma Colon
## 1  0.00      4.35      4.35      4.35      0.00 4.35      4.35      0.00 0.00 0.00
## 2  7.14      0.00     14.29      0.00      0.00 7.14     21.43      7.14 0.00 0.00
## 3  5.66      5.66      5.66      3.77      1.89 3.77     20.75      3.77 3.77 0.00
## 4  0.00      0.00      0.00      0.00      0.00 0.00     12.50      0.00 12.50 0.00
## 5  0.00     13.73      0.00      5.88      0.00 0.00     33.33      7.84 7.84 5.88
## 6  0.00      0.00      0.00      0.00      0.00 0.00     10.00      0.00 10.00 0.00
## 7  1.82      0.00      5.45      0.00      0.00 0.00     30.91      5.45 3.64 0.00
## 8  0.00     13.79      0.00     10.34      0.00 0.00     17.24      6.90 10.34 0.00
## 9  0.00      0.00      0.00      0.00      0.00 0.00     14.29      0.00 0.00 0.00
## 10 0.00      0.00      5.00      0.00      0.00 0.00      5.00      0.00 5.00 0.00
## 11 0.00      6.25      6.25      6.25      0.00 0.00     18.75      6.25 0.00 0.00
## 12 0.00      5.77      3.85      3.85      0.00 0.00     25.00      5.77 0.00 0.00
## 13 0.00      3.57      7.14      0.00      0.00 7.14     14.29      3.57 0.00 7.14
## 14 5.26     10.53      5.26      5.26      5.26 5.26     15.79     15.79 0.00 0.00
## 15 2.08      4.17      4.17      2.08      0.00 2.08     31.25     10.42 6.25 0.00
##      SemiC QMark Exclam Dash Quote Apostro Parenth OtherP
## 1  0.00 0.00      0.00 0.00 4.35      4.35      0      4.35
## 2  0.00 0.00      0.00 0.00 14.29      0.00      0     21.43
## 3  0.00 0.00      0.00 0.00 13.21      5.66      0     20.75
## 4  0.00 0.00      0.00 0.00 0.00      0.00      0     12.50
## 5  3.92 0.00      0.00 1.96 3.92      3.92      0     31.37
## 6  0.00 0.00      0.00 0.00 0.00      0.00      0     10.00
## 7  3.64 0.00      1.82 0.00 3.64      0.00      0     30.91
## 8  0.00 0.00      0.00 0.00 0.00      0.00      0     17.24
## 9  0.00 14.29      0.00 0.00 0.00      0.00      0     14.29
## 10 0.00 0.00      0.00 0.00 0.00      0.00      0      5.00
## 11 0.00 0.00      0.00 0.00 6.25      6.25      0     18.75
## 12 0.00 1.92      0.00 0.00 1.92      1.92      0     17.31
## 13 0.00 0.00      0.00 0.00 0.00      0.00      0     14.29
## 14 0.00 0.00      0.00 0.00 0.00      0.00      0     15.79
## 15 2.08 0.00      0.00 0.00 4.17      0.00      0     31.25
```

```
head(cap_sentimentGI, 15)
```

```
##      docname Segment WC      WPS Sixltr   Dic positive negative AllPunc Period
## 1   text1      1 23 23.000000 21.74 13.04      4.35      8.70      4.35      0.00
## 2   text2      2 14 5.500000 21.43 0.00      0.00      0.00     21.43      7.14
## 3   text3      3 53 22.500000 16.98 15.09      7.55      7.55     20.75      3.77
## 4   text4      4 8 7.000000 12.50 12.50      0.00     12.50     12.50      0.00
## 5   text5      5 51 9.250000 27.45 13.73      1.96     11.76     33.33      7.84
## 6   text6      6 10 9.000000 30.00 10.00     10.00      0.00     10.00      0.00
## 7   text7      7 55 14.333333 20.00 7.27      7.27      0.00     30.91      5.45
## 8   text8      8 29 12.000000 34.48 10.34      3.45      6.90     17.24      6.90
## 9   text9      9 7 6.000000 14.29 0.00      0.00      0.00     14.29      0.00
## 10 text10     10 20 19.000000 15.00 30.00     15.00     15.00      5.00      0.00
## 11 text11     11 16 7.000000 18.75 6.25      0.00      6.25     18.75      6.25
## 12 text12     12 52 9.000000 23.08 13.46      5.77      7.69     25.00      5.77
## 13 text13     13 28 12.500000 25.00 3.57      3.57      0.00     14.29      3.57
## 14 text14     14 19 5.333333 31.58 5.26      5.26      0.00     15.79     15.79
## 15 text15     15 48 12.000000 33.33 6.25      2.08      4.17     31.25     10.42
##      Comma Colon SemiC QMark Exclam Dash Quote Apostro Parenth OtherP
```




```
## 1  0.00  0.00  0.00  0.00  0.00  0.00  4.35  4.35  0  4.35
## 2  0.00  0.00  0.00  0.00  0.00  0.00 14.29  0.00  0 21.43
## 3  3.77  0.00  0.00  0.00  0.00  0.00 13.21  5.66  0 20.75
## 4 12.50  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0 12.50
## 5  7.84  5.88  3.92  0.00  0.00  1.96  3.92  3.92  0 31.37
## 6 10.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0 10.00
## 7  3.64  0.00  3.64  0.00  1.82  0.00  3.64  0.00  0 30.91
## 8 10.34  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0 17.24
## 9  0.00  0.00  0.00 14.29  0.00  0.00  0.00  0.00  0 14.29
##10  5.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0  5.00
##11  0.00  0.00  0.00  0.00  0.00  0.00  6.25  6.25  0 18.75
##12  0.00  0.00  0.00  1.92  0.00  0.00  1.92  1.92  0 17.31
##13  0.00  7.14  0.00  0.00  0.00  0.00  0.00  0.00  0 14.29
##14  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0.00  0 15.79
##15  6.25  0.00  2.08  0.00  0.00  0.00  4.17  0.00  0 31.25
```

Both produce dataframe objects that can later be manipulated to conduct more in-depth analysis. Let's see how "anger" and "joy" language varies across groups and time using the NRC dictionary.

```
cap_sentiment_df <- cbind.data.frame(cap_sentimentNRC,data_capitalism)
cap_sentiment_df <- cap_sentiment_df[cap_sentiment_df$date_created > "2019-11-19",]


ggplot(cap_sentiment_df, aes(x=date_created, y=anger, color = mem_name))+
  stat_smooth() +
  labs(title="Anger in Capitalism", x="Date", y = "Anger", las=2)+
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

```
## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```



```
ggplot(cap_sentiment_df, aes(x=date_created, y=joy, color = mem_name))+
  stat_smooth() +
  labs(title="Joy in Capitalism", x="Date", y = "Joy", las=2)+
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

```
## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```



Sentiment Analysis-2.bb

One of the limitations of sentiment analysis is the variation we get on similar sentiments when using different dictionaries. For example, analyzing *polarity* in our data (positive language minus negative language), we obtain different results from GI and NRC.

```
polarity_NRC <- cap_sentimentNRC$positive -
cap_sentimentNRC$negative

polarity_GI <- cap_sentimentGI$positive -
cap_sentimentGI$negative

cor(polarity_NRC,polarity_GI)
```

```
## [1] 0.5070789
```

As a general recommendation, it is always good practice to explore the dictionary we are going to use before actually using it. This is particularly important when working with text in another language. Most dictionaries (commercial or otherwise) that have been evaluated and validated, have been tested and evaluated in English. If using another language, it might be a good idea to check how words are classified. Another common practice is to use Google Translate API or Microsoft Translator API on the original corpus. This can be costly depending on the size of the corpus and the results, like with any other text-analysis tool, need to be constantly validated (Lucas et al. 2015).

**Table 1: Matrix  $D$** 

	$W_1$	$W_2$	$W_3$	$W_m$
$D_1$	3	4	0	1
$D_2$	0	3	0	4
$D_3$	1	2	2	0
$D_m$	1	0	2	3

Figure 2: Matrix D

## 9 Analyzing a corpus: Topic models and document distance

Topic models were primarily developed to summarize unstructured text, use words within documents to infer their topic, and as a form of dimension reduction. It allows social scientists to use topics as a form of measurement, as we are often interested in how observed covariates drive trends in language. For example,

- Political attention patterns in Senate floor speeches (Quinn et al. 2010)
- Attention senators allocate to press releases (Grimmer 2010).
- Climate change “skepticism” in reports and communications by think tanks and interest groups (Bou-saills and Coan 2016).

Topic models are “unsupervised” methods. As such, they require a great deal of human supervision, especially when it comes to validating the results.

### 9.1 Document-term Matrix

Topic models are a broad class of Bayesian generative models that encode problem-specific structure into an estimation of categories (Grimmer and Stewart 2013; Blei et al. 2010). Statistically, a topic is a probability mass function over words. The idea of topic models is that each document exhibits a topic in some proportion: each document is a distribution over topics, and each topic is a distribution over words.

We can take our corpus and turn it into a matrix that reflects this concept. We call it a document-term matrix (*dtm*).

- A corpus of  $n$  documents  $D_1, D_2, D_3 \dots D_n$
- Vocabulary of  $m$  words  $W_1, W_2, W_3 \dots W_m$
- The value of  $i, j$  cell gives the frequency count of word  $W_j$  in document  $D_i$

Latent Dirichlet Allocation (LDA) converts the *dtm* into two lower-dimension matrices,  $M_1$  and  $M_2$ :

- $M_1$  is a  $N \times K$  document-topic matrix
- $M_2$  is a  $K \times M$  topic-term matrix

**Table 2:** Matrix  $M_1$

	$K_1$	$K_2$	$K_3$	$K_p$
$D_1$	1	0	0	1
$D_2$	1	1	0	0
$D_3$	1	0	1	1
$D_m$	1	0	0	0

Figure 3: Matrix M1

**Table 3:** Matrix  $M_2$

	$W_1$	$W_2$	$W_3$	$W_m$
$K_1$	1	0	1	1
$K_2$	0	1	0	0
$K_3$	1	1	1	0
$K_p$	1	0	0	0

Figure 4: Matrix M2

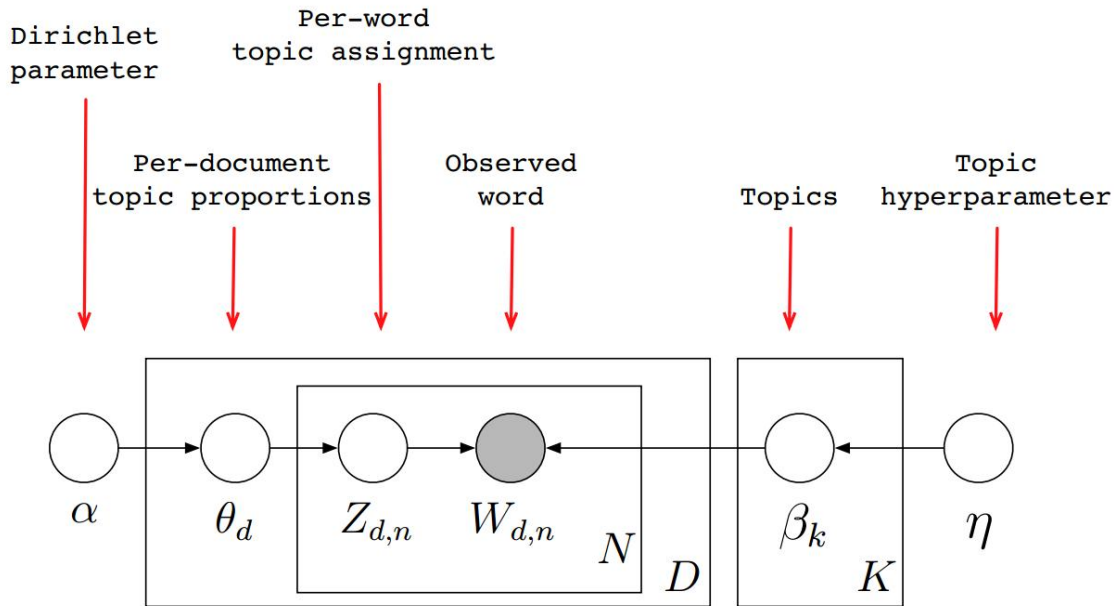


Figure 5: Plate notation of Latent Dirichlet Allocation

LDA estimates the distribution over words for each topic, and the proportion of a topic in each document. You can read about the math behind the two-step process in Grimmer and Stewart (2013).

- $\alpha$   $\rightarrow$  document-topic density: higher  $\alpha$  means documents contain more topics, lower  $\alpha$  means a document contains fewer topics.
- $\beta$   $\rightarrow$  topic-word density. higher  $\beta$  means topics have more words, lower  $\beta$  means topics have fewer words.
- Number of topics  $\rightarrow$  this is specified in advance, or can be chosen to optimize model fit (we will get back to this point). The “statistically optimal” topic count is usually too high for the topics to be interpretable or useful.

## 9.2 Structural Topic Models (STM)

To apply (and visualize) LDA, we are going to be using an extension of LDA known as Structural Topic Models (STM).

$$STM = LDA + Metadata$$

STM provides two ways to include contextual information to “guide” the estimation of the model. First, topic prevalence can vary by metadata (e.g. Republicans talk about military issues more than Democrats). Second, topic content can vary by metadata (e.g. Republicans talk about military issues differently from Democrats).

We can run STM using the `stm` package. The `stm` package includes the complete workflow (i.e. from raw text to figures), and if you are planning to use it in the future I highly encourage you to check this and this. `stm()` takes our `dfm` and produces topics. If we do not specify any prevalence terms, then it will estimate an LDA. Since this is a Bayesian approach, it is recommended you set a seed value for future replication. We also need to set  $K$  number of topics. How many topics are the right number of topics? There is no good number. Too many pre-specified topics and the categories might be meaningless. Too few, and you might be

piling together two or more topics. Note that changes to a) the number of topics, b) the prevalence term, c) the omitted words, d) the seed value, can (greatly) change the outcome. Here is where validation becomes crucial (for a review see Wilkerson and Casas 2017).

Using our dataset, I will use `stm` to estimate the topics surrounding “capitalism” on Twitter. As my prevalence term, I add the position of each authority. I set my number of topics at 10 (but with a corpus this big I should probably set it at ~30 and work my way up from there).

```
dfm_toks_stem <- dfm(cap_toks_stem)

# After trimming, some documents were left with no tokens so I will eliminate those before running my model
dfm_toks_sub <- dfm_subset(dfm_toks_stem, ntoken(dfm_toks_stem) > 0)

cap_TM <- stm(dfm_toks_sub, K = 10, seed = 1984,
              prevalence = ~left_right,
              init.type = "Spectral")
```

The nice thing about the `stm()` function is that it allows us to see in “real-time” what is going on within the black box. We can summarize the process in the following way (this is similar to a collapsed Gibbs sampling, which the `stm()` function sort of uses):

1. Go through each document, and randomly assign each word in the document to one of the topics  $t \in k$ .
2. Notice that this random assignment already gives both topic representations of all the documents and word distributions of all the topics (albeit not very good ones).
3. So to improve on them, for each document  $W$  do the following:
  - 3.1 Go through each word  $w$  in  $W$ 
    - 3.1.1 And for each topic  $t$ , compute two things:
      - 3.1.1.1  $p(t|W)$  = the proportion of words in document  $W$  that are currently assigned to topic  $t$ , and
      - 3.1.1.2  $p(w|t)$  = the proportion of assignments to topic  $t$  over all documents that come from this word  $w$ .
    - Reassign  $w$  a new topic, where we choose topic  $t$  with probability  $p(t|W) * p(w|t)$ . It is worth noting that according to our generative model, this is essentially the probability that topic  $t$  generated word  $w$ , so it makes sense that we resample the current word’s topic with this probability. (Also, I’m glossing over a couple of things here, in particular the use of priors/pseudocounts in these probabilities.)
    - 3.1.1.3 In other words, in this step, we’re assuming that all topic assignments except for the current word in question are correct, and then updating the assignment of the current word using our model of how documents are generated.
4. After repeating the previous step a large number of times, you’ll eventually reach a roughly steady state where your assignments are pretty good. So use these assignments to estimate the topic mixtures of each document (by counting the proportion of words assigned to each topic within that document) and the words associated with each topic (by counting the proportion of words assigned to each topic overall).

(This explanation was taken from here). Let’s explore the topics produced:

```
labelTopics(cap_TM)

## Topic 1 Top Words:
## Highest Prob: peopl, like, good, end, never, look, realli
## FREX: like, radic, gift, art, vegan, digit, empire
## Lift: 80s, aral, chichisup, coffin, dutch, etcetera, feet
## Score: like, peopl, gift, look, good, diana, princess
## Topic 2 Top Words:
## Highest Prob: social, markruffalo, rich, million, kill, mark, socialist
## FREX: markruffalo, million, mark, ruffalo, worth, millionair, hollywood
## Lift: deepstat, 30m, breitbartnew, bridgetphetasi, clue, cronyism, darkwatersmovi
```

```

##      Score: markruffalo, mark, ruffalo, million, kill, worth, social
## Topic 3 Top Words:
##      Highest Prob: just, say, fuck, know, problem, much, american
##      FREX: fuck, black, love, leftist, friday, u, sex
##      Lift: asshol, bent, blk, bourgeois, chart, custom, dick
##      Score: fuck, just, hate, black, friday, love, problem
## Topic 4 Top Words:
##      Highest Prob: social, market, free, state, anti, climat, chang
##      FREX: climat, global, differ, u.s, competit, challeng, threat
##      Lift: , 150th, 17th, 1990s, 1a2a, 4human, 4sale
##      Score: market, state, climat, free, social, global, stage
## Topic 5 Top Words:
##      Highest Prob: get, money, go, now, give, everyon, busi
##      FREX: money, ethic, pretend, speak, disgust, get, sleep
##      Lift: agreed, artworkte, barackobama, ben, bigot, blatant, brianstelt
##      Score: get, money, go, give, ethic, everyon, bad
## Topic 6 Top Words:
##      Highest Prob: system, need, human, world, profit, live, exploit
##      FREX: imperi, resourc, individu, relationship, growth, longer, altern
##      Lift: agricultur, anarchist_black, arab, autonomi, bake, bolivian, cathol
##      Score: system, human, profit, resourc, exploit, product, natur
## Topic 7 Top Words:
##      Highest Prob: work, can, us, worker, time, class, better
##      FREX: ewarren, debt, labour, insur, hour, corrupt, care
##      Lift: catastrophe, davi, ewarren, insecur, mitig, patient, unafford
##      Score: us, work, can, corrupt, pay, warren, poor
## Topic 8 Top Words:
##      Highest Prob: freedom, berni, public, polici, protest, capit, failur
##      FREX: berni, polici, protest, failur, sander, sensand, leecamp
##      Lift: pete, _michelangelo_, (, , , 11warrior, 13_moth
##      Score: zerogbadillion, jimmy_dor, ajc4other, alllibertynew, blysx, cptseammonkey, decakarjeffre
## Topic 9 Top Words:
##      Highest Prob: make, think, right, even, labor, part, instead
##      FREX: instead, fuel, imposs, consequ, content, vision, cheap
##      Lift: , , , , analyz, beckett, behaviour
##      Score: make, think, labor, even, right, instead, el
## Topic 10 Top Words:
##      Highest Prob: thing, way, also, creat, power, societi, great
##      FREX: thing, seem, issu, anandwrit, hierarchi, essay, racism
##      Lift: acceleration, ambit, assert, blackston, calebmaupin, copyright, demis
##      Score: thing, also, creat, racism, societi, way, power


```

FREX weights words by their overall frequency and how exclusive they are to the topic. Lift weights words by dividing by their frequency in other topics, therefore giving higher weight to words that appear less frequently in other topics. Similar to lift, score divides the log frequency of the word in the topic by the log frequency of the word in other topics (Roberts et al. 2013). Bischof and Airoldi (2012) show the value of using FREX over the other measures.

You can use the `plot()` function to show the topics.

```
plot(cap_TM, type = "summary", labeltype = "frex") # or prob, lift score
```

## Top Topics



Topics-1.bb

### Expected Topic Proportions

If you want to see a sample of a specific topic:


```
findThoughts(cap_TM, texts = texts(cap_corp)[docnames(dfm_toks_sub)], topics = 4)
```

```
##
## Topic 4:
## "Neoliberalism is the 20th-century resurgence of 19th-century ideas associated with laissez-fa
##
## Not that vague Nate
## @LuckyHeronSay @LibDems The LibD`s,Tories & anybody who supports capitalism & the cuts o
## Only full blooded socialism,the truth--& an IMPLEMENTED socialist alt. will suffice 2 beat them.
## Blairites have IMPLEMENTED Tory cuts etc.
## Lab could be in trouble in their areas.
## @PHPatriot_1 @zoothorn69 The connection between #Christianity, #capitalism, our #Constitution, a
```

We can (should/must) run some diagnostics. There are two qualities that we are looking for in our model: semantic coherence and exclusivity. Exclusivity is based on the FREX labeling matrix. Semantic coherence is a criterion developed by Mimno et al. (2011) and it maximizes when the most probable words in a given topic frequently co-occur together. Mimno et al. (2011) show that the metric correlates well with human judgement of topic quality. Yet, it is fairly easy to obtain high semantic coherence so it is important to see it in tandem with exclusivity. Let's see how exclusive are the words in each topic:


```
dotchart(exclusivity(cap_TM), labels = 1:10)
```





We can also see the semantic coherence of our topics –words a topic generates should co-occur often in the same document–:

```
dotchart(semanticCoherence(cap_TM,dfm_toks_sub), labels = 1:10)
```




Coherence-1.bb

We can also see the overall quality of our topic model:

```
topicQuality(cap_TM,dfm_toks_sub)
```

```
## [1] -153.0452 -112.0533 -156.6526 -152.9175 -144.6928 -148.9168 -149.8053
## [8] -222.4208 -187.0373 -155.3585
## [1] 9.617573 9.436403 9.562088 9.482883 9.343743 9.494762 9.505163 9.616828
## [9] 9.345646 9.572110
```



On their own, both metrics are not really useful (what do those numbers even mean?). They are useful when we are looking for the “optimal” number of topics.

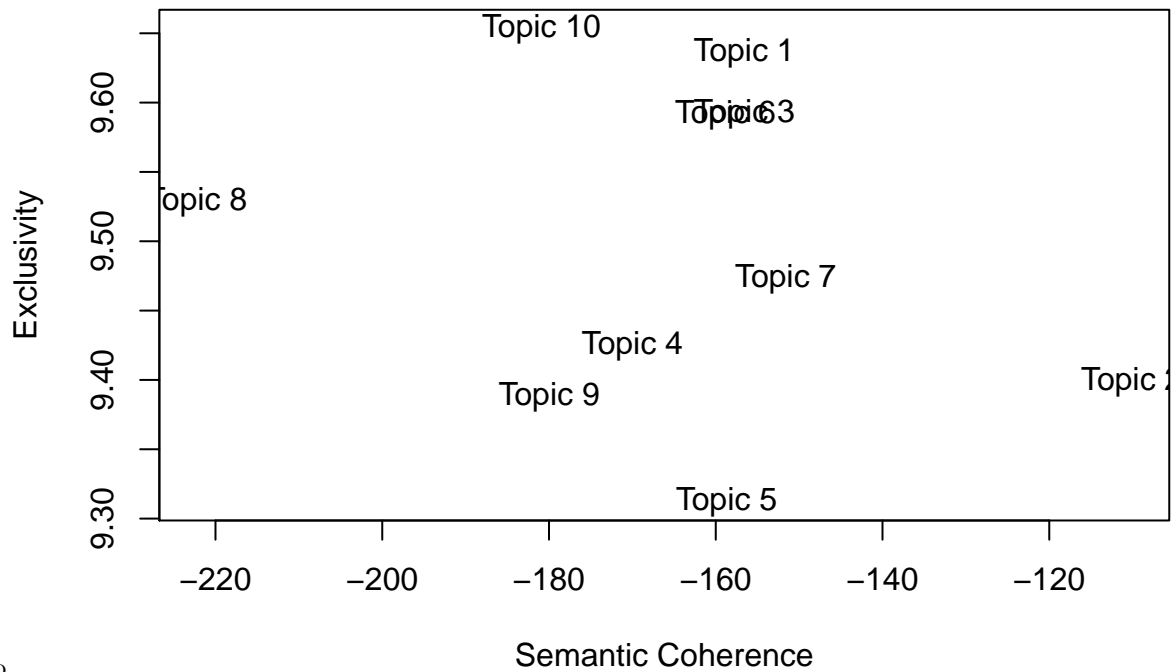
```
cap_TM_0 <- manyTopics(dfm_toks_sub,
  prevalence = ~ left_right,
  K = c(10,15,20), runs=2,
  max.em.its = 50,
  init.type = "Spectral") # It takes around 250 iterations for the model to converge
```

We can now compare the performance of each model based on their semantic coherence and exclusivity:

```
k_10 <- cap_TM_0$out[[1]] # k_10 is an stm object which can be explored and used like any other topic model
k_15 <- cap_TM_0$out[[2]]
k_20 <- cap_TM_0$out[[3]]


# I will just graph the 'quality' of each model:
topicQuality(k_10,dfm_toks_sub)
```

```
## [1] -156.5871 -110.0929 -156.5661 -169.9799 -158.6913 -158.8270 -151.5906
## [8] -222.2478 -179.9702 -180.8846
## [1] 9.635785 9.398349 9.591756 9.424861 9.312281 9.591102 9.473077 9.528538
## [9] 9.387953 9.653311
```



```
topicQuality(k_15,dfm_toks_sub)
```


```
## [1] -176.3001 -117.4834 -184.9537 -158.2778 -169.1092 -216.5443 -166.8171
## [8] -188.7352 -235.7356 -141.5557 -200.8848 -176.1947 -154.4283 -170.0446
## [15] -183.5428
## [1] 9.752277 9.752145 9.806760 9.624878 9.615049 9.353169 9.652194 9.655456
## [9] 9.689886 9.686355 9.597133 9.759681 9.661525 9.349216 9.976027
```



```
topicQuality(k_20,dfm_toks_sub)
```

```
## [1] -155.8810 -165.9696 -207.8772 -142.5555 -106.7070 -157.1547 -153.8948
## [8] -169.3887 -155.9115 -213.4163 -169.2706 -200.3757 -175.9227 -168.2318
```

```
## [15] -171.0044 -249.4002 -159.4595 -190.0719 -183.0467 -256.6249
## [1] 9.751484 9.536017 9.698744 9.651605 9.721448 9.893554 9.533039 9.723660
## [9] 9.769429 9.422858 9.648541 9.854435 9.718662 9.553076 9.872522 9.895744
## [17] 9.851317 9.818935 9.816059 9.865039
```



by K-3.bb

Maybe we have some theory about the difference in topic prevalence across sides (or across parties). We can see the topic proportions in our topic model object:

```
head(cap_TM$theta)
```

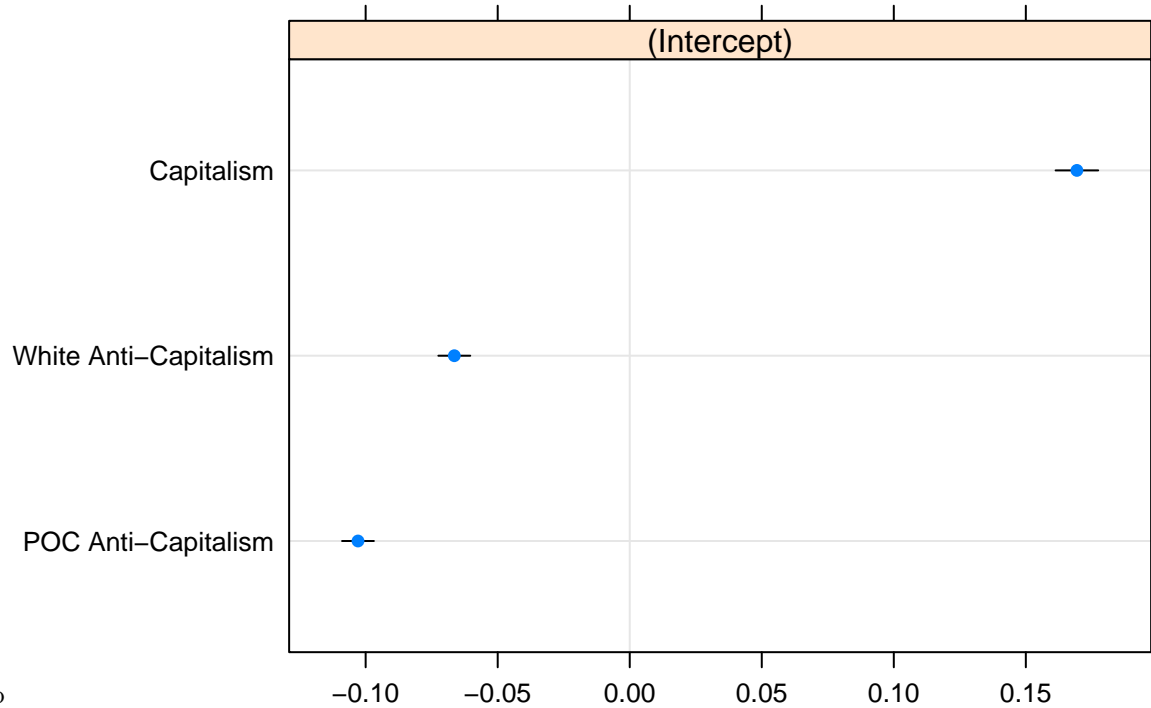
```
##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
## [1,] 0.07542875 0.02024515 0.49208606 0.07776768 0.09624133 0.06542261
## [2,] 0.15249491 0.18153916 0.08170145 0.03862463 0.30355151 0.05864056
## [3,] 0.16350298 0.13006020 0.07387521 0.07534854 0.12696537 0.10668717
## [4,] 0.31965533 0.03596207 0.14965279 0.06928192 0.09203642 0.09834901
## [5,] 0.02665616 0.01267243 0.03409487 0.50179732 0.01986142 0.23534687
## [6,] 0.02036968 0.01553662 0.01370014 0.04234740 0.01344656 0.02801444
##           [,7]      [,8]      [,9]     [,10]
## [1,] 0.06764590 0.006605026 0.05305056 0.04550695
## [2,] 0.09278333 0.007145652 0.03322783 0.05029096
## [3,] 0.16127230 0.009931679 0.06931750 0.08303905
## [4,] 0.08599592 0.009958856 0.05526251 0.08384517
## [5,] 0.05048596 0.015786213 0.01677402 0.08652474
## [6,] 0.03254694 0.797484935 0.01129259 0.02526068
```

What about connecting this info to our dfm and see if there are differences in the proportion topic 6 is addressed by each side.

```
cap_prev <- data.frame(topic2 = cap_TM$theta[,2], docvars(dfm_toks_sub))
lmer_topic2 <- lmer(topic2 ~ (1 | left_right), data = cap_prev)
dotplot(ranef(lmer_topic2, condVar = TRUE))
```

```
## $left_right
```

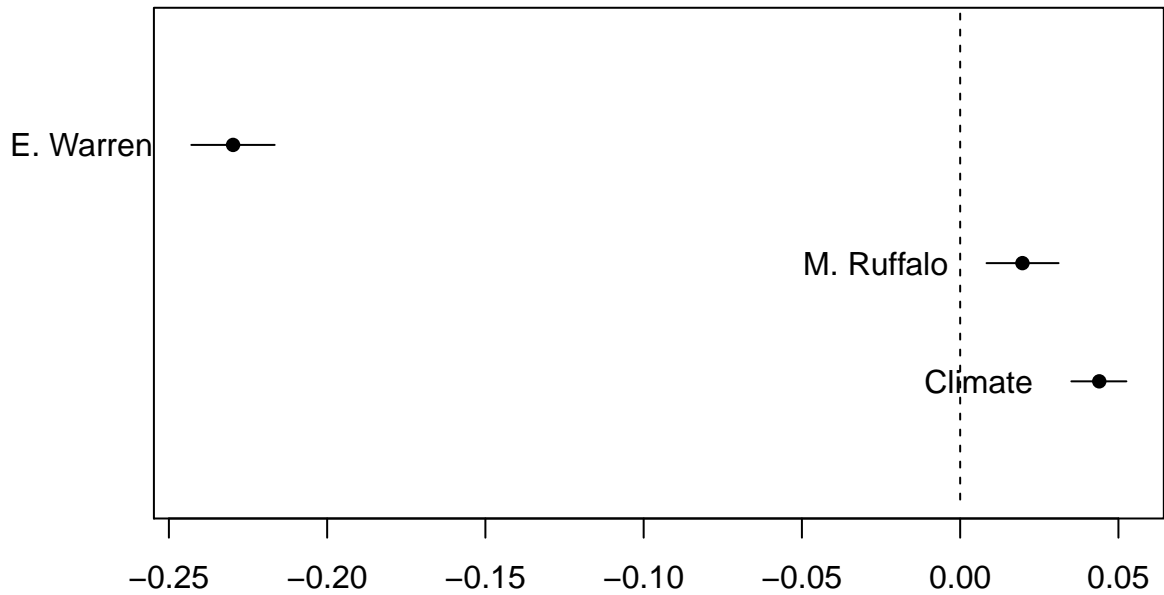
## left\_right



Topics-1.bb

Makes sense: anti-capitalist bunch together and away from capitalists. We can do something similar with the `stm` function directly. We just need to specify the functional form and add the document variables.

```
cap_topics <- estimateEffect(c(2,4,7) ~ left_right_dum, cap_TM, docvars(dfm_toks_sub)) # You can compare
plot(cap_topics, "left_right_dum", method = "difference",
     cov.value1 = "Anti-Capitalism",
     cov.value2 = "Capitalism",
     labeltype = "custom",
     xlab = "More Liberal ... More Conservative",
     custom.labels = c('E. Warren', 'M. Ruffalo', 'Climate'),
     model = cap_TM)
```



Effect-1.bb

More Liberal ... More Conservative

## 10 Analyzing a corpus: Scaling models

So far, we have explored tools that provide information about the text. We can also use the text to obtain information about the authors. The **Wordfish** model developed by Slapin and Proksch (2008), for example, positions the authors of documents in an (ideological) scale. How? In politics, the frequency with which politician  $i$  uses word  $k$  is drawn from a Poisson distribution:

$$w_{ik} \sim \text{Poisson}(\lambda_{ik})$$

$$\lambda_{ik} = \exp(\alpha_i + \beta_k + \gamma_i \beta_k)$$

with latent parameters:

- $\alpha_i$  is the “loquaciousness” of politician  $i$
- $\beta_k$  is the frequency of word  $k$
- $\gamma_k$  is the discrimination parameter of word  $k$
- $\gamma_i$  is the politician’s ideological position

The parameters of interest are the  $\gamma_i$ ’s, the position of the parties in each election year, and the  $\beta_k$ ’s because they allow us to analyze which words differentiate between party positions.

The main assumption is that, indeed,  $\lambda_{ik}$  is generated by the parameters previously described. Let’s believe for a second that the peer-review system works and use the `textmodel_wordfish()` function to estimate the positions of our authorities in our corpus.

```
## I will subset my data since as it is there are too many documents:
```

```
data_capitalism <- data_capitalism %>%
  group_by(nameauth) %>%
  mutate(count_auth = n())

data_capitalism_sub <- data_capitalism[data_capitalism$count_auth>8,]
```

```

## I will also concatenate all the tweets by author:
data_capitalism_sub <- data_capitalism_sub %>%
  group_by(nameauth) %>%
  mutate(text_conca = paste0(text_clean, collapse = " "))

## And finally drop dups:
data_capitalism_sub <- data_capitalism_sub[!duplicated(data_capitalism_sub$nameauth),]

cap_corp_sub <- corpus(data_capitalism_sub$text_conca,
  docvars = data.frame(author = data_capitalism_sub$nameauth,
    left_right = data_capitalism_sub$mem_name
  ))

## Again my dfm
cap_dmf_sub <- dfm(cap_corp_sub,
  remove_punct = TRUE,
  remove_numbers = TRUE,
  remove = stopwords("english"),
  stem = T)

cap_wfish <- textmodel_wordfish(cap_dmf_sub, dir = c(1,5)) #Does not really matter what the starting va
summary(cap_wfish)

```

```

##
## Call:
## textmodel_wordfish.dfm(x = cap_dmf_sub, dir = c(1, 5))
##
## Estimated Document Positions:
##      theta      se
## text1  0.27208 0.06189
## text2 -0.99178 0.04885
## text3 -0.77402 0.04343
## text4  0.37626 0.06202
## text5  0.53266 0.05240
## text6  0.54913 0.06721
## text7 -1.70036 0.01783
## text8 -1.04105 0.05488
## text9  1.18058 0.04790
## text10 0.46719 0.10257
## text11 -0.99258 0.06303
## text12 0.92237 0.07797
## text13 0.07457 0.07749
## text14 -0.72464 0.07788
## text15 -0.70685 0.07443
## text16 -0.45435 0.11727
## text17 0.39764 0.09232
## text18 0.26899 0.08314
## text19 -1.07466 0.06112
## text20 -0.10482 0.24448
## text21 1.61314 0.03896
## text22 0.15266 0.05828
## text23 0.30705 0.11761
## text24 1.34803 0.03473

```

```
## text25  1.36736 0.05965
## text26  2.14261 0.02236
## text27 -1.08596 0.04948
## text28 -0.67270 0.04066
## text29 -0.86851 0.06740
## text30  1.00146 0.04266
## text31 -1.78151 0.01650
##
## Estimated Feature Scores:
##      trump conserv republican democrat    amp liber    link common    bond
## beta  0.1831  0.3257    -0.1235   0.3558 0.7277 -1.428  0.8157 -0.1832  0.3446
## psi  -1.2239 -1.7671    -3.0482  -1.7695 0.6283 -2.368 -3.1269 -2.1448 -2.3282
##      capit imperi coloni   white supremaci patriarchi jingoism    allow
## beta 0.09024 -0.9719 -0.4575  0.4393   -0.1455   -0.1455   0.5496 -0.1235
## psi  2.40763 -2.0284 -2.7410 -1.6457   -3.0527   -3.0527  -3.7443 -3.0482
##      maintain    relat luxuri lifestyl  expens  global    major planet    thus
## beta  -0.2717 -0.3965  0.1206    1.413 -0.1662  0.5085  0.01487  1.195  0.626
## psi  -3.0839 -2.0252 -2.6122   -2.523 -2.3640 -0.8460 -0.94706 -2.051 -2.376
##      call realli personifi virtual
## beta  0.5985  0.3681   -0.3714 -0.3714
## psi  -1.3573 -0.9440   -3.1151 -3.1151
```


This is an interesting exercise, since the reasoning behind **wordfish** is similar to the one behind network analysis. In network analysis, rather than looking at the text, we look at the connections made from Tweets and ReTweets. Let's see how both approaches compare:

```
cap_preds <- predict(cap_wfish, interval = "confidence")
cap_pos <- data.frame(docvars(cap_corp_sub),
                      cap_preds$fit) %>%
  arrange(fit)

cap_pos <- cap_pos[order(cap_pos$fit),] # sort

ggplot(cap_pos, aes(x = fit, y = 1:nrow(cap_pos), xmin = lwr, xmax = upr,color = left_right)) +
  geom_point() +
  geom_errorbarh(height = 0) +
  scale_y_continuous(labels = cap_pos$docs, breaks = 1:nrow(cap_pos)) +
  labs(x = "Position", y = "User") +
  ggtitle("Estimated Positions")
```





Positions from WordFish-1.bb

External validation, beibi! (?)

We can also turn around the scaling and see where each word is positioned on the same left-right scale as the authors. Here is the “Eiffel Tower” Slapin and Proksch (2008) and of scaled words:

```
wscores <- data.frame(word = cap_wfish$features,
  score = cap_wfish$beta,
  offset = cap_wfish$psi)

wscores <- wscores[wscores$score<5,]

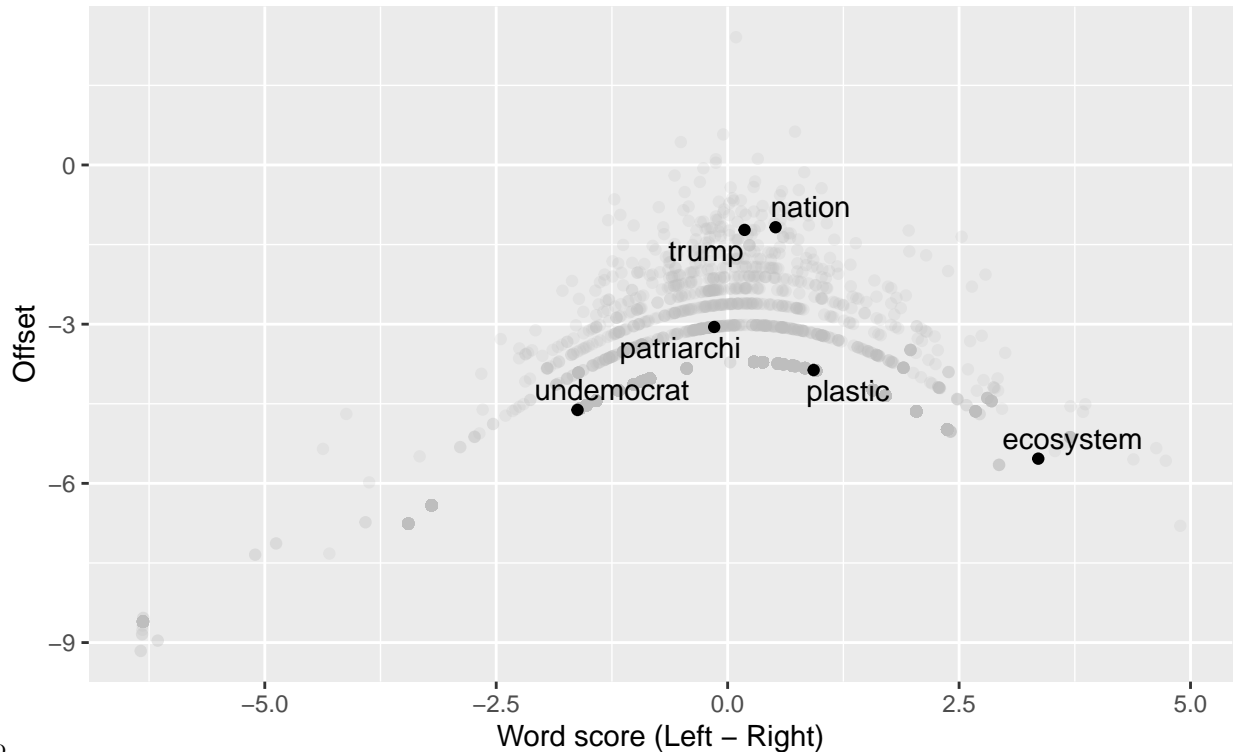
testwords <- c("oligarch", "ecosystem", "undemocrat","plastic",
  "nation", "trump", "patriarchi")

testscores <- wscores %>%
  filter(word %in% testwords) %>%
  arrange(score)

ggplot(wscores, aes(score, offset, label = word)) +
  geom_point(color = "grey", alpha = 0.2) +
  geom_text_repel(data = testscores, col = "black") +
  geom_point(data = testscores) +
  labs(x = "Word score (Left - Right)", y = "Offset") +
  ggtitle("Estimated position of words",
    subtitle = "Nota: Parameter offset is proportional to the frequency of each word.")
```

## Estimated position of words

Nota: Parameter offset is proportional to the frequency of each word.



Tower-1.bb

One important limitation of **wordfish** is that it assumes that all the documents are addressing the same topic, which is not necessarily the case. But there are scaling models for every taste (for example this and this), so this should not be too much of a problem. Still, it is incumbent upon the researcher to choose the model that best reflects the data and the needs.

## 11 Analyzing a corpus: Natural Language Processing (NLP)

Working with natural language is not a solved problem. Language is messy and ever-changing and evolving. It takes us the better part of our childhood to learn it, “it is hard for the scientist who attempts to model the relevant phenomena, and it is hard for the engineer who attempts to build systems that deal with natural language input or output.” (Kornai 2008))

“Statistical NLP aims to do statistical inference for the field of natural language. Statistical inference in general consists of taking some data (generated in accordance with some unknown probability distribution) and then making some inference about this distribution.” (Manning and Schutze 1999)

NLP has been used for a while in software: word predictors in your phone, spell-checkers, spam filtering, etc. Its implementation in political science is still limited but the possibilities are vast. NLP makes it possible to move beyond simply establishing connections to investigating the state of relationships, from example by moving from ‘whom’ to ‘who did what to whom.’ (Welbers et al. 2017)

We will be testing three advanced NLP techniques: lemmatization, part-of-speech (POS) tagging, and dependency parsing.

### 11.0.1 Lemmatization

Much like stemming, but rather than cutting off words a dictionary is used to replace terms with their **lemmas**. More accurate at normalizing words with different verb forms (e.g. “gave” and “give”), which is a desirable quality when pre-processing a corpus.

### 11.0.2 Part-of-speech tagging

Syntactic categories for words, such as nouns, verbs, articles, and adjectives.

From Welbers et al. 2017: “This information can be used to focus an analysis on certain types of grammar categories, for example, using nouns and proper names to measure similar events in news items (Welbers, et al., 2016), or using adjectives to focus on subjective language (De Smedt and Daelemans, 2012).”

```
cap_parsed <- spacy_parse(data_capitalism$text_clean)
```

```
## Finding a python executable with spaCy installed...
```

```
## spaCy (language model: en) is installed in /usr/local/bin/python3
```

```
## successfully initialized (spaCy Version: 2.0.16, language model: en)
```

```
## (python options: type = "python_executable", value = "/usr/local/bin/python3")
```

```
head(cap_parsed,25)
```

##	doc_id	sentence_id	token_id	token	lemma	pos	entity
## 1	text1	1	1	if	if	ADP	
## 2	text1	1	2	you	-PRON-	PRON	
## 3	text1	1	3	're	be	VERB	
## 4	text1	1	4	having	have	VERB	
## 5	text1	1	5	trouble	trouble	NOUN	
## 6	text1	1	6	understanding	understand	VERB	
## 7	text1	1	7	capitalism	capitalism	NOUN	
## 8	text1	1	8	just	just	ADV	
## 9	text1	1	9	remember	remember	VERB	
## 10	text1	1	10	there	there	ADV	
## 11	text1	1	11	are	be	VERB	
## 12	text1	1	12	only	only	ADV	
## 13	text1	1	13	4	4	NUM	CARDINAL_B
## 14	text1	1	14	websites	website	NOUN	
## 15	text1	1	15	you	-PRON-	PRON	
## 16	text1	1	16	use	use	VERB	
## 17	text1	1	17	in	in	ADP	
## 18	text1	1	18	2019	2019	NUM	DATE_B
## 19	text1	2	1	and	and	CCONJ	
## 20	text1	2	2	you	-PRON-	PRON	
## 21	text1	2	3	hate	hate	VERB	
## 22	text1	2	4	all	all	DET	
## 23	text1	2	5	of	of	ADP	
## 24	text1	2	6	them	-PRON-	PRON	
## 25	text2	1	1	"	"	PUNCT	

### 11.0.3 Dependency parsing

Dependency parsing provides the syntactic relations between tokens. For example, “Kendrick” is related to “Lamar”, thus recognizing “Kendrick Lamar” as a single entity. This can be particularly useful if you are searching for certain types of entities (e.g. locations, institutions, etc.) in a corpus. We can see the differences in the location mentioned by capitalists and anti-capitalists:

```
cap_entities <- entity_extract(cap_parsed)
head(cap_entities, 20)
```

##	doc_id	sentence_id	entity	entity_type
## 1	text6	1	Ariana_Grande_Breaks	PERSON
## 2	text6	2	Bernie_Sanders	PERSON
## 3	text7	3	\n	GPE
## 4	text7	4	Bernie	PERSON
## 5	text7	4	America	GPE
## 6	text7	5	\n	GPE
## 7	text7	6	@cornelWest	ORG
## 8	text7	6	@chaunceydevega	ORG
## 9	text7	6	\n	GPE
## 10	text10	1	Arundhati_Roy	PERSON
## 11	text11	2	Elites_Lost_Their_Grip	PERSON
## 12	text11	2	Time	ORG
## 13	text12	2	Dem	NORP
## 14	text12	3	\n\n	ORG
## 15	text12	3	\n\n_Dem	PERSON
## 16	text12	4	\n\n_Capitalism	PERSON
## 17	text12	5	\n\n	ORG
## 18	text12	5	\n\n_Capitalism	PERSON
## 19	text14	2	American	NORP
## 20	text15	1	DemocraticSocialism	ORG

```
# Similar to extract but converts multi-word entities into single "tokens":
# cap_entities <- entity_consolidate(cap_parsed_dep)
```

```
cap_persons <- cap_entities$entity[cap_entities$entity_type == "PERSON"]
cap_persons <- unique(cap_persons)
head(cap_persons, 40)
```

## [1]	"Ariana_Grande_Breaks"	"Bernie_Sanders"	"Bernie"
## [4]	"Arundhati_Roy"	"Elites_Lost_Their_Grip"	"\n\n_Dem"
## [7]	"\n\n_Capitalism"	"Ilhan_Omar"	"Free_Speech"
## [10]	"Mark_Ruffalo"	"Ariana_Grande"	"John_Legend"
## [13]	"Chrissy_Teigen"	"Elizabeth_Warren"	"Harry_Haywood"
## [16]	"Angela_Davis"	"Mark"	"We_Need_'_Revolution_'"
## [19]	"Einstein"	"Wayne"	"Nonstop"
## [22]	"Deval_Patrick_'s"	"Hitler"	"\n\n_Me"
## [25]	"\n\n_Marx"	"Danny"	"Greed"
## [28]	"Deval_Patrick_'s"	"Jobs"	"Colin"
## [31]	"Nigel_Farage"	"Lamar"	"Pete"
## [34]	"Rhodes_Scholar"	"Obama"	"\n\n_Spreading"
## [37]	"Black"	"Hillary_Clinton"	"John_Cornyn"
## [40]	"Lech_Walesa"		

In other contexts, spaCyR can recognize which subject is doing the action and which subject is on the receiving end. Van Atteveldt et al. (2017) use this to analyze who is attacking whom in news about the

Gaza war. Here we can see something similar at work. Let's....

```
cap_parsed_dep <- spacy_parse(data_capitalism$text_clean, dependency = TRUE, entity = TRUE, lemma = FALSE)
head(cap_parsed_dep, 25)
```

##	doc_id	sentence_id	token_id	token	pos	tag	head_token_id	dep_rel
## 1	text1	1	1	if	ADP	IN	4	mark
## 2	text1	1	2	you	PRON	PRP	4	nsubj
## 3	text1	1	3	're	VERB	VBP	4	aux
## 4	text1	1	4	having	VERB	VBG	9	advcl
## 5	text1	1	5	trouble	NOUN	NN	4	dobj
## 6	text1	1	6	understanding	VERB	VBG	5	acl
## 7	text1	1	7	capitalism	NOUN	NN	6	dobj
## 8	text1	1	8	just	ADV	RB	9	advmod
## 9	text1	1	9	remember	VERB	VBP	9	ROOT
## 10	text1	1	10	there	ADV	EX	11	expl
## 11	text1	1	11	are	VERB	VBP	9	ccomp
## 12	text1	1	12	only	ADV	RB	13	advmod
## 13	text1	1	13	4	NUM	CD	14	nummod
## 14	text1	1	14	websites	NOUN	NNS	11	attr
## 15	text1	1	15	you	PRON	PRP	16	nsubj
## 16	text1	1	16	use	VERB	VBP	14	relcl
## 17	text1	1	17	in	ADP	IN	11	prep
## 18	text1	1	18	2019	NUM	CD	17	pobj
## 19	text1	2	1	and	CCONJ	CC	3	cc
## 20	text1	2	2	you	PRON	PRP	3	nsubj
## 21	text1	2	3	hate	VERB	VBP	3	ROOT
## 22	text1	2	4	all	DET	DT	3	dobj
## 23	text1	2	5	of	ADP	IN	4	prep
## 24	text1	2	6	them	PRON	PRP	5	pobj
## 25	text2	1	1	"	PUNCT	``	2	amod
##	entity							
## 1								
## 2								
## 3								
## 4								
## 5								
## 6								
## 7								
## 8								
## 9								
## 10								
## 11								
## 12								
## 13	CARDINAL_B							
## 14								
## 15								
## 16								
## 17								
## 18	DATE_B							
## 19								
## 20								
## 21								
## 22								
## 23								

```
## 24
## 25
```

spacyR can also detect other attributes of tokens in a text:

```
cap_parsed_att <- spacy_parse(data_capitalism$text,
  additional_attributes = c("like_num", "like_url"),
  lemma = FALSE, pos = FALSE, entity = FALSE)
head(cap_parsed_att,30)
```

##	doc_id	sentence_id	token_id	token	like_num	like_url
## 1	text1	1	1	if	FALSE	FALSE
## 2	text1	1	2	you	FALSE	FALSE
## 3	text1	1	3	're	FALSE	FALSE
## 4	text1	1	4	having	FALSE	FALSE
## 5	text1	1	5	trouble	FALSE	FALSE
## 6	text1	1	6	understanding	FALSE	FALSE
## 7	text1	1	7	capitalism	FALSE	FALSE
## 8	text1	1	8	just	FALSE	FALSE
## 9	text1	1	9	remember	FALSE	FALSE
## 10	text1	1	10	there	FALSE	FALSE
## 11	text1	1	11	are	FALSE	FALSE
## 12	text1	1	12	only	FALSE	FALSE
## 13	text1	1	13	4	TRUE	FALSE
## 14	text1	1	14	websites	FALSE	FALSE
## 15	text1	1	15	you	FALSE	FALSE
## 16	text1	1	16	use	FALSE	FALSE
## 17	text1	1	17	in	FALSE	FALSE
## 18	text1	1	18	2019	TRUE	FALSE
## 19	text1	2	1	and	FALSE	FALSE
## 20	text1	2	2	you	FALSE	FALSE
## 21	text1	2	3	hate	FALSE	FALSE
## 22	text1	2	4	all	FALSE	FALSE
## 23	text1	2	5	of	FALSE	FALSE
## 24	text1	2	6	them	FALSE	FALSE
## 25	text2	1	1	"	FALSE	FALSE
## 26	text2	1	2	capitalism	FALSE	FALSE
## 27	text2	1	3	not	FALSE	FALSE
## 28	text2	1	4	cronyism	FALSE	FALSE
## 29	text2	1	5	"	FALSE	FALSE
## 30	text2	1	6	poster	FALSE	FALSE

## 12 Other techniques not covered

### 12.1 Word networks

How are words and authors connected? We can create networks of authors as nodes and edges based on the overlap in words between authors. Let's use the `textnet` package created by Chris Bail:

```
# library(devtools)
# install_github("cbail/textnets")
library(textnets)
```

```
## Loading required package: udpipe
```

```

## Warning: package 'udpipe' was built under R version 3.5.2
## Loading required package: ggraph
## Warning: package 'ggraph' was built under R version 3.5.2
## Loading required package: networkD3
## Warning: replacing previous import 'dplyr::union' by 'igraph::union' when
## loading 'textnets'
## Warning: replacing previous import 'dplyr::as_data_frame' by
## 'igraph::as_data_frame' when loading 'textnets'
## Warning: replacing previous import 'dplyr::groups' by 'igraph::groups' when
## loading 'textnets'
# I will subset my data but lower the threshold to get more authorities:

data_capitalism_sub <- data_capitalism[data_capitalism$count_auth>4,]

## I will also concatenate all the tweets by author:
data_capitalism_sub <- data_capitalism_sub %>%
  group_by(nameauth) %>%
  mutate(text_conca = paste0(text_clean, collapse = " "))

## And finally drop dups:
data_capitalism_sub <- data_capitalism_sub[!duplicated(data_capitalism_sub$nameauth),]

# We are only going to be using nouns.
data_capitalism_sub$text_conca_nocap <- str_remove_all(data_capitalism_sub$text_conca, "[Cc]apitalism")

# We prep our data (takes a while):
prepped_cap <- PrepText(data_capitalism_sub, groupvar = "nameauth", textvar = "text_conca_nocap", nodevar = "nameauth")

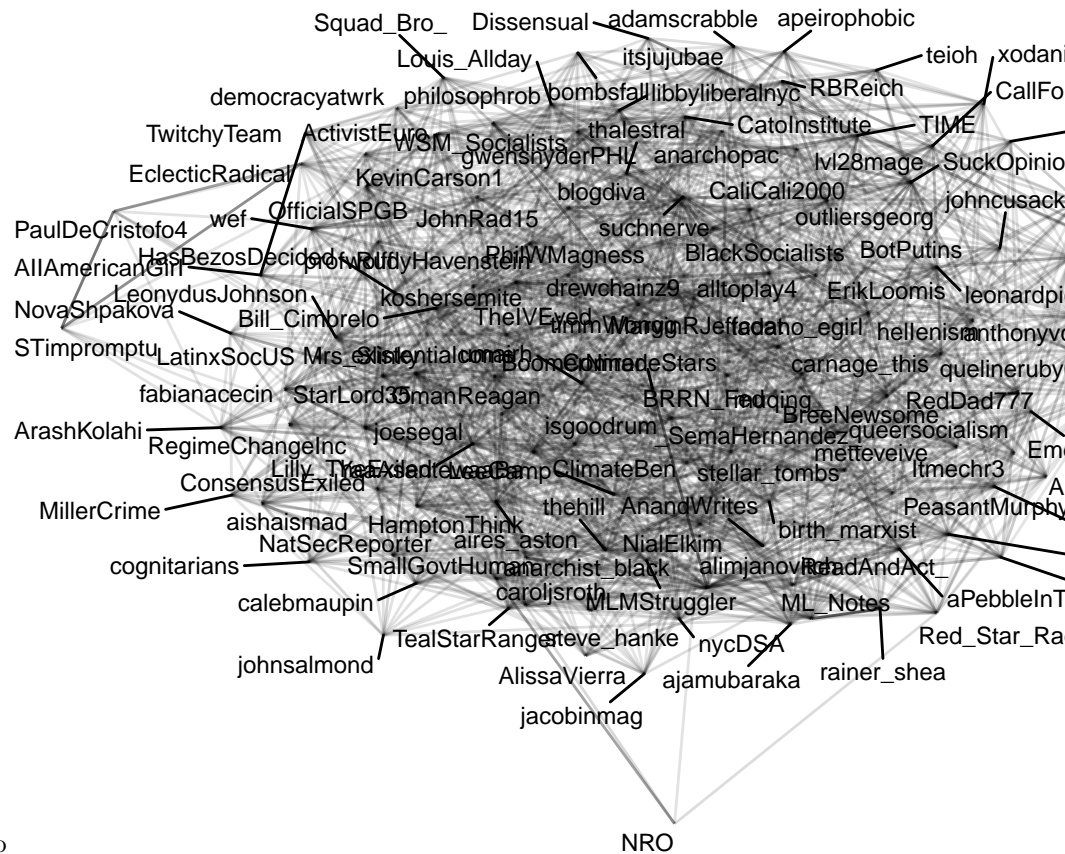
## Downloading udpipe model from https://raw.githubusercontent.com/jwijffels/udpipe.models.ud.2.4/master/udpipe.models.ud.2.4.tar.gz
## Visit https://github.com/jwijffels/udpipe.models.ud.2.4 for model license details

To create the adjacency matrix for the network, we use the CreateTextnet() function. The cells of the adjacency matrix are the transposed crossproduct of the term-frequency inverse-document frequency (TFIDF) for overlapping terms between two documents for PrepText and the matrix product of TFIDF crossproduct (See Bail, 2016).

cap_text_network <- CreateTextnet(prepped_cap)
VisTextNet(cap_text_network, label_degree_cut = 0)

## Using `stress` as default layout

```



and Visualizing the Network-1.bb


A mess... We can see the communities and compare with the communities that we got from the original network analysis.

```
cap_communities <- TextCommunities(cap_text_network)

cap_communities_full <- cbind.data.frame(cap_communities, data_capitalism_sub$to_membership)
colnames(cap_communities_full)[3] <- "membership_net"
cap_communities_full$modularity_class <- as.numeric(cap_communities_full$modularity_class)

ggplot(cap_communities_full, aes(x=membership_net, y=modularity_class)) +
  geom_point() +
  geom_jitter(width = 0.2, height = 0.2)
```





Communities-1.bb

We would expect a bit more exclusivity of membership, but oh well...

## 12.2 Cosine-similarity

How similar are two texts? If we think of each text as a vector in space, we can think of the angle between the two vectors as their “distance”. The smaller the angle  $\theta$ , the closer and the more similar. If we had a measure of each text (e.g. TF-IDF) we could compute the cosine of the angle between them. How? Math. We must solve the dot product for the  $\cos \theta$ :

$$\vec{a} \cdot \vec{b} = \|\vec{a}\| \|\vec{b}\| \cos \theta$$

$$\cos \theta = \frac{\vec{a} \cdot \vec{b}}{\|\vec{a}\| \|\vec{b}\|}$$

That is the cosine similarity formula. Cosine similarity will generate a metric that says how related are two documents by looking at the angle instead of the magnitude:

## 12.3 Word embeddings

Word embeddings use a similar intuition as cosine similarities. By using vector representation of text, we can estimate how much alike two words are rather than treating words as single independent units (which is the case for dictionaries).


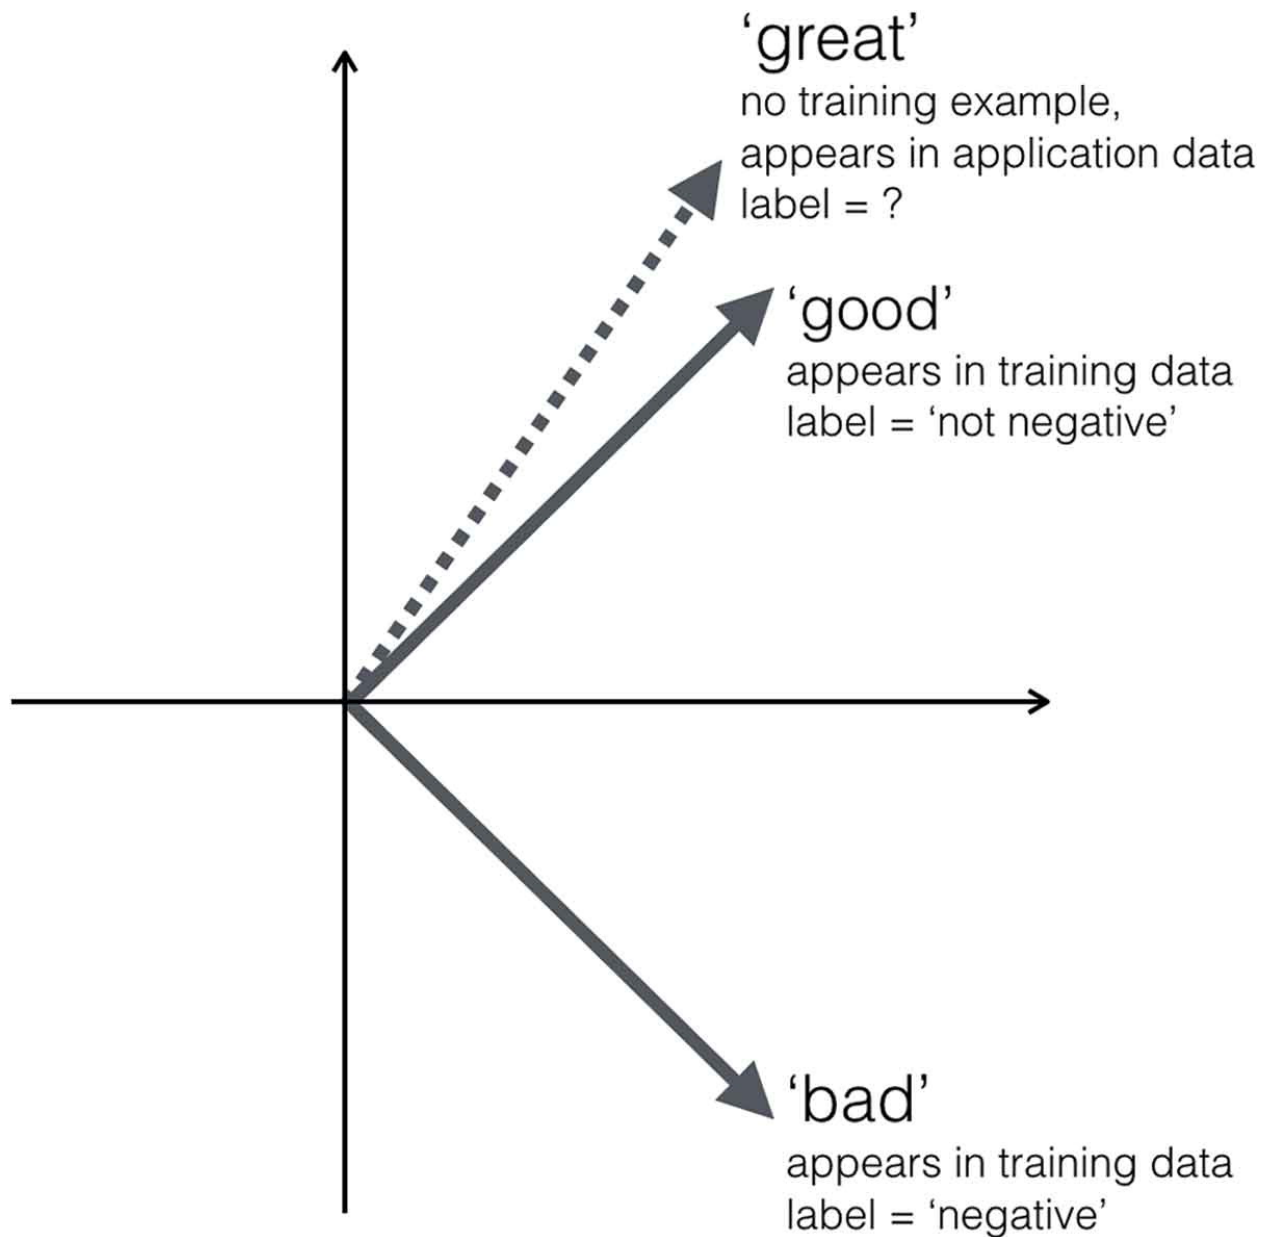


Figure 6: Cosine similarity



We can train our models to predict how words are related. This requires a previously annotated training data set and some technical skills (and the pre-processing of the data). The literature suggest that word embeddings can produce better results than similar “bag-of-words” approaches. It is particularly attractive since these models can be ran in most languages, helping us overcome the limitations from canned English-focused dictionaries.

## 13 Goodbye Notes

- If your are interested in text analysis, the internet is your best friend. I am also happy to help. And to be your friend.
- The resources suggested at the beginning will help you navigate text analysis and see how it fits within your research.
- Learn to scrape the interwebs.
- Learn regex.
- Ask.
- Remember: 2 locations, 3 versions for all your data. Backups must be *Regular, Automatic, and Incremental*.